

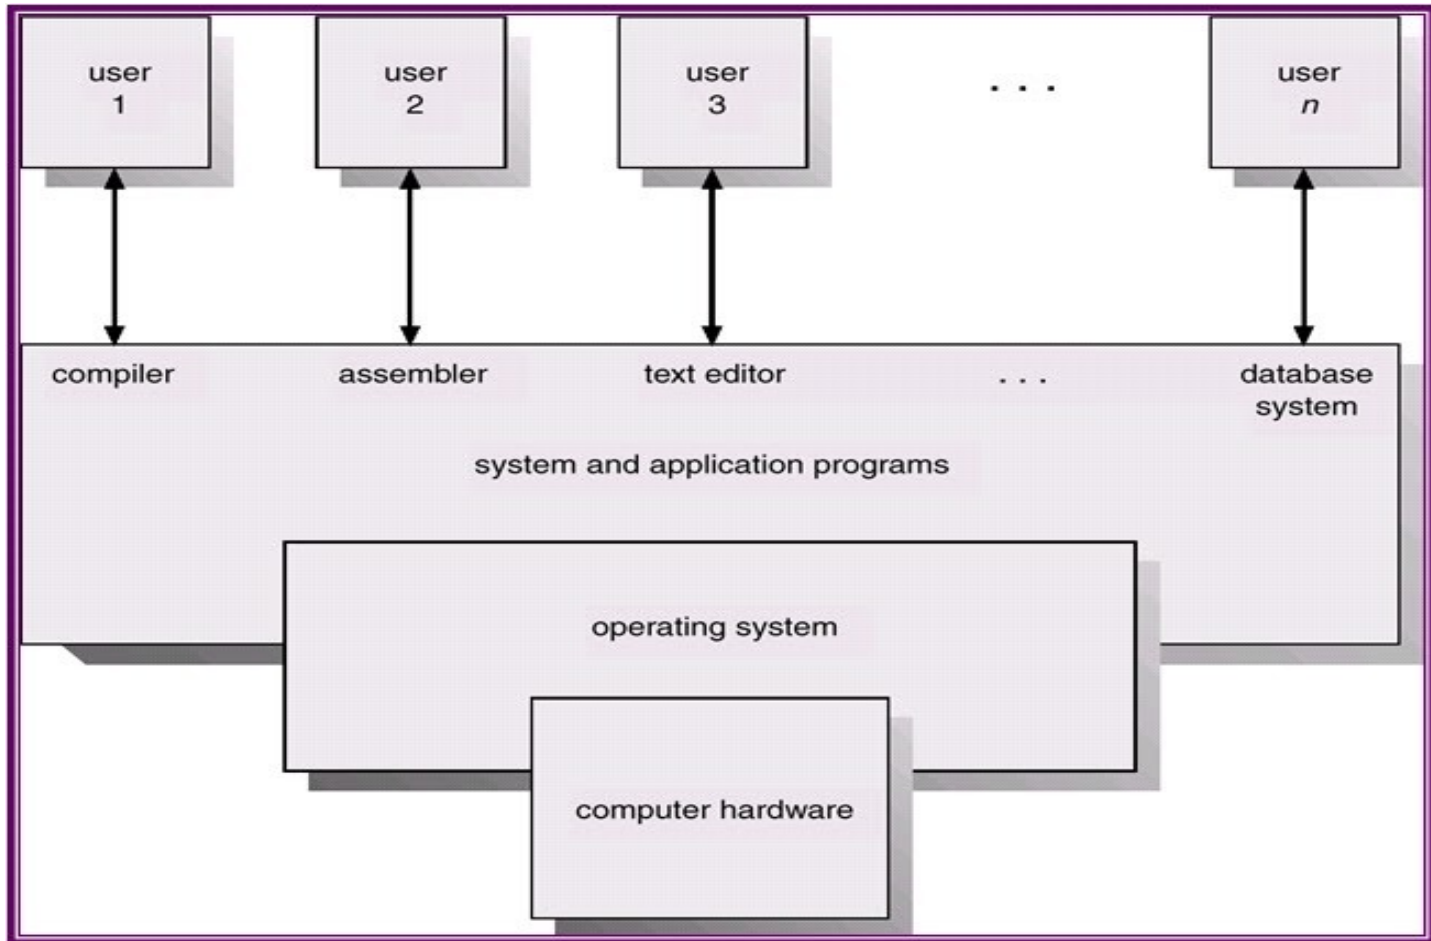


UNIVERSITY OF  
BIRMINGHAM

# Computer Systems Concurrency & Threads



# A view of an operating System



# What does this imply?

- ◆ Multiple 'users'
  - ◆ Working independently
  - ◆ The operating system:
    - ◆ Protects memory from interference from other processes
    - ◆ Interleaves execution of each process to:
      - ◆ Maximise resource utilisation
      - ◆ Responsiveness of each process
  - ◆ Preserves the states of each process so it can resume execution



# Historically ....

- ◆ Computers were expensive:

- ◆ Maximise utilisation

- ◆ Now:

- ◆ Users expect concurrent execution of multiple tasks

- ◆ Stream music

- ◆ Update software

- ◆ Edit file

- ◆ Test program

- ◆ Monitor SNS (Social Networking Services)

- ◆ etc.



# So?

- ◆ The problem is that these processes do not *always* work independently:
  - ◆ They can compete for access to resources:
    - ◆ Devices, files, data etc. that are local (or remote)
  - ◆ They can co-operate – through messages, shared memory, files etc.
  - ◆ Applications are distributed:
    - ◆ So, potentially, millions of people are reading or updating information *at the same time*
    - ◆ SNS, bank accounts, on line shopping ....



# And ...

- ◆ We need to build applications that can interleave multiple independent (but potentially conflicting) tasks:
  - ◆ Consider an app with a GUI:
    - ◆ do we want it to freeze when it is performing some time consuming task?
  - ◆ What about a database of bank accounts?
    - ◆ Do we want to wait in a queue to perform a transaction?
  - ◆ What about programmers' efficiency & effectiveness?
    - ◆ Do we want them to have to 'build an OS' for every app?
- ◆ What about the effectiveness of use of our computing resources?
  - ◆ We want to maximise the efficiency

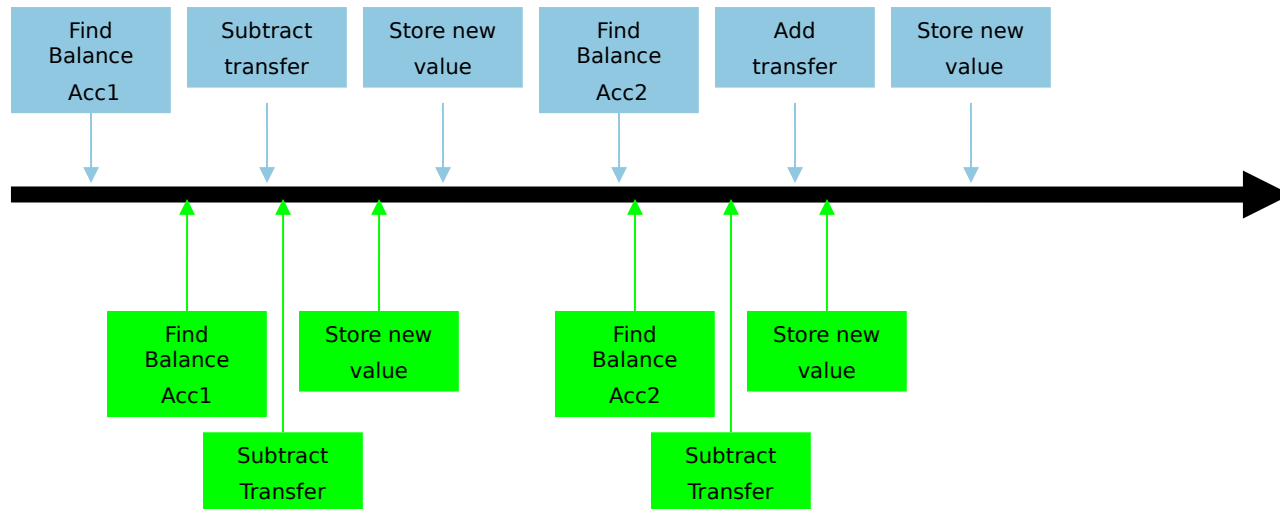


This leads us to two things:

- ◆ How can we manage concurrency?
  - ◆ How can we ensure that our concurrent execution behaves correctly?
  - ◆ How can we do this conveniently?
  - ◆ How can we do this efficiently?
    - ◆ To maintain responsiveness
- ◆ How can we provide programming mechanisms *within a program* to allow concurrency?
  - ◆ Threads



# Let's think about a bank transfer



This is unlikely:

- but possible!
- and hard to debug!





# Concurrency

- ◆ Introduction to Concurrency
- ◆ OS Concerns and Process Interaction
- ◆ Process Synchronization
- ◆ Critical Section Problem
- ◆ Requirements for Solutions to CS Problem
- ◆ Software Solutions
- ◆ Hardware Solutions
- ◆ High Level OS Support

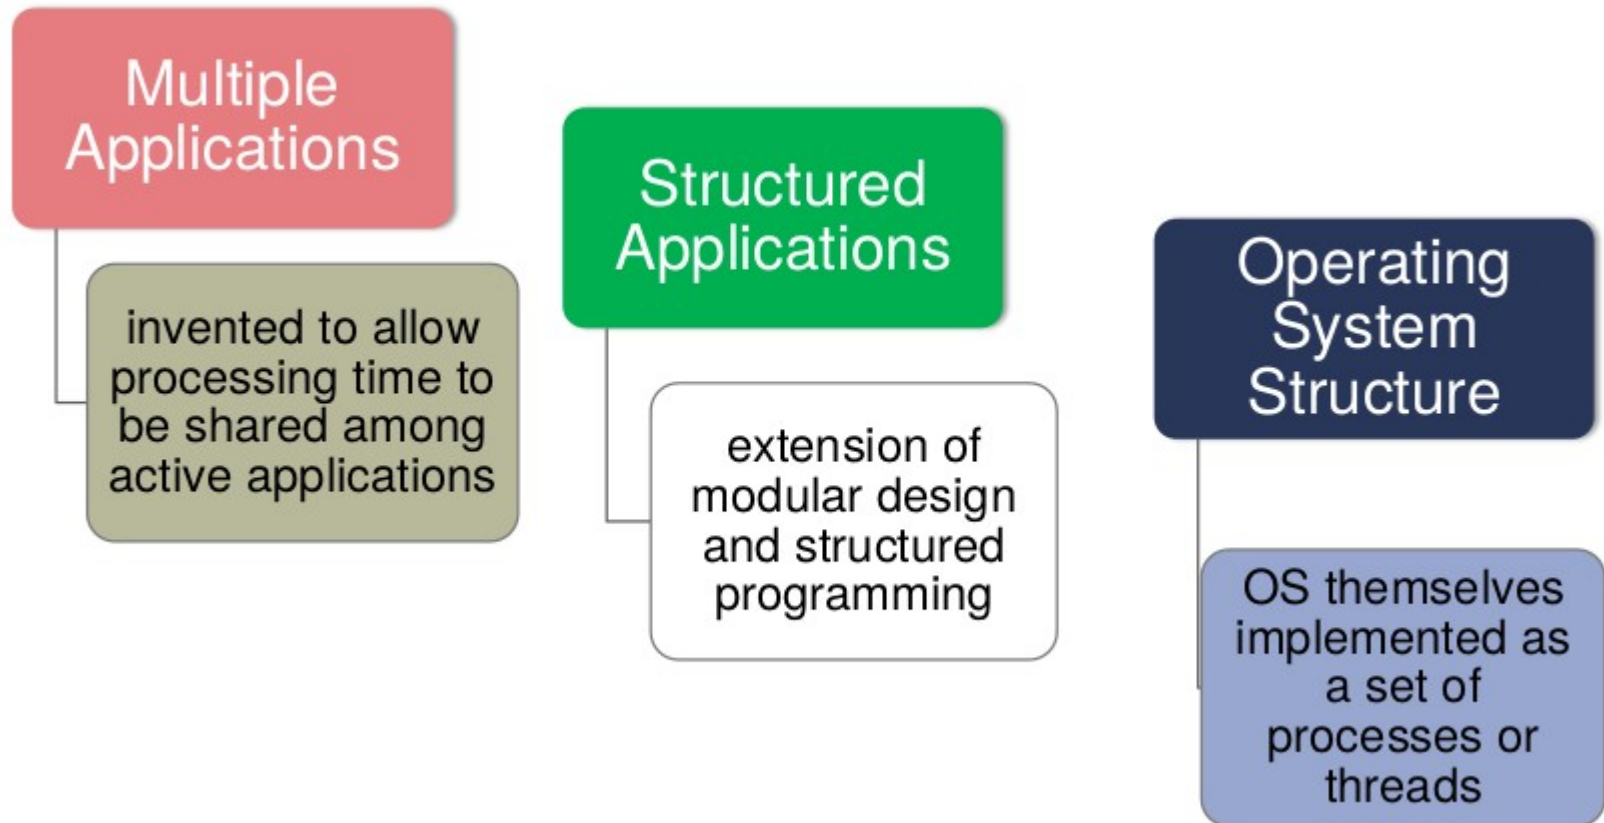


# Multiple Processes

- ◆ Operating System design is concerned with the management of **processes** and **threads**
  - Multiprogramming (Multitasking)
  - Multiprocessing
  - Distributed Processing

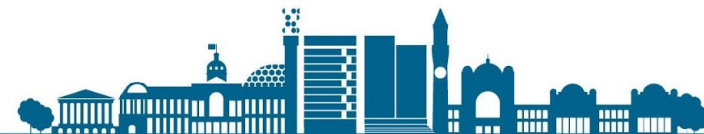


# Concurrency in Different Contexts



# Concurrency: Key Terms

<b>atomic operation</b>	A function or action implemented as a sequence of one or more instructions that appears to be indivisible; that is, no other process can see an intermediate state or interrupt the operation. The sequence of instruction is guaranteed to execute as a group, or not execute at all, having no visible effect on system state. Atomicity guarantees isolation from concurrent processes.
<b>critical section</b>	A section of code within a process that requires access to shared resources and that must not be executed while another process is in a corresponding section of code.
<b>deadlock</b>	A situation in which two or more processes are unable to proceed because each is waiting for one of the others to do something.
<b>livelock</b>	A situation in which two or more processes continuously change their states in response to changes in the other process(es) without doing any useful work.
<b>mutual exclusion</b>	The requirement that when one process is in a critical section that accesses shared resources, no other process may be in a critical section that accesses any of those shared resources.
<b>race condition</b>	A situation in which multiple threads or processes read and write a shared data item and the final result depends on the relative timing of their execution.
<b>starvation</b>	A situation in which a runnable process is overlooked indefinitely by the scheduler; although it is able to proceed, it is never chosen.



# Principles of Concurrency

## Interleaving and Overlapping

- ◆ Can be viewed as examples of **concurrent** processing
- ◆ Both present the **same** problems

**Uniprocessor** – the relative timing of execution of processes cannot be predicted:

- ◆ Depends on activities of other processes
- ◆ The way the OS handles interrupts
- ◆ Scheduling policies of the OS



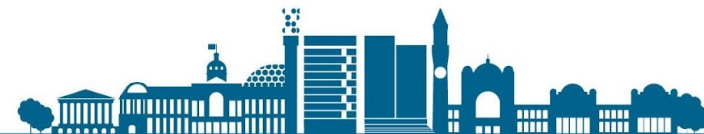
# Difficulties of Concurrency

- ◆ Sharing of global resources
- ◆ Difficult for the OS to manage the allocation of resources optimally
- ◆ Coordination of 'parallel' and asynchronous 'processes'
  - ◆ Ensuring correct behaviour
  - ◆ Debugging of programming errors:
    - ◆ behaviour can be *non-deterministic* and, therefore, hard to identify and reproduce



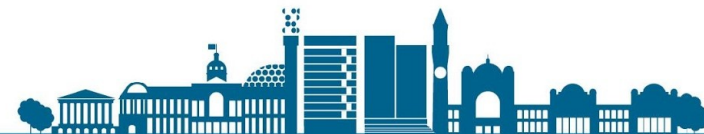
# Operating System Concerns

- ◆ Design and management issues raised by the existence of concurrency:
- ◆ The OS must:
  - be able to **keep track** of various **processes**
  - **allocate** and **de-allocate** resources for each active process
  - **protect** the data and physical resources of each process against *interference by other processes*
  - ensure that the processes and outputs are **independent** of the **processing speed**



# Process Interaction

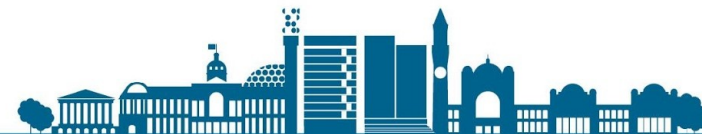
Degree of Awareness	Relationship	Influence that One Process Has on the Other	Potential Control Problems
Processes unaware of each other	Competition	<ul style="list-style-type: none"><li>• Results of one process independent of the action of others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Mutual exclusion</li><li>• Deadlock (renewable resource)</li><li>• Starvation</li></ul>
Processes indirectly aware of each other (e.g., shared object)	Cooperation by sharing	<ul style="list-style-type: none"><li>• Results of one process may depend on information obtained from others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Mutual exclusion</li><li>• Deadlock (renewable resource)</li><li>• Starvation</li><li>• Data coherence</li></ul>
Processes directly aware of each other (have communication primitives available to them)	Cooperation by communication	<ul style="list-style-type: none"><li>• Results of one process may depend on information obtained from others</li><li>• Timing of process may be affected</li></ul>	<ul style="list-style-type: none"><li>• Deadlock (consumable resource)</li><li>• Starvation</li></ul>






# Resource Competition

- ◆ Concurrent processes come into conflict when they are competing for use of the same resource
  - for example: I/O devices, memory, processor time, clock
- ◆ In the case of competing processes **three control problems** must be faced:
  - Mutual Exclusion – to ensure correct behaviour
  - Deadlocks
  - Starvation

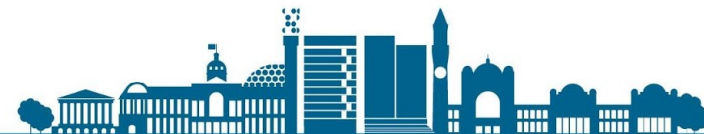


# Synchronization of Processes – Example

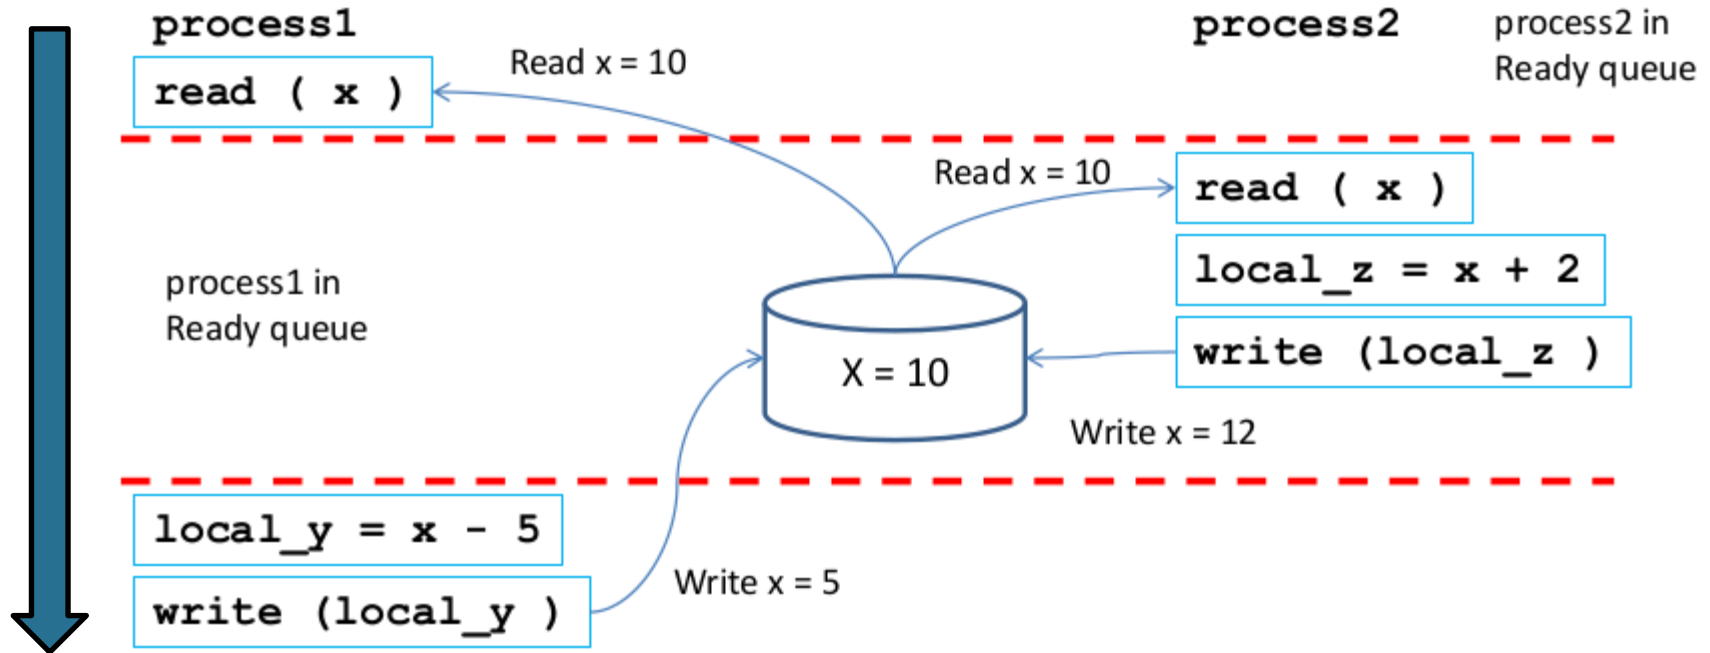


## ◆ We expect

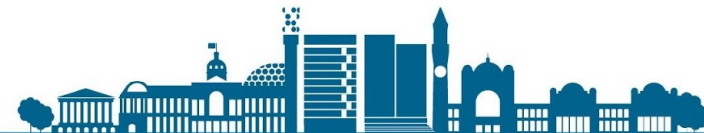
- When process1 finishes, shared variable  $x$  is *reduced by 5*
- When process2 finishes, shared variable  $x$  is *increased by 2*



# Synchronization of Processes – Example




- ◆ Context switches may occur *at any time*
- ◆ **Process 2 has its result overwritten by process 1**
- ◆ **Process 1 operates with outdated information**





# What is a Race Condition?

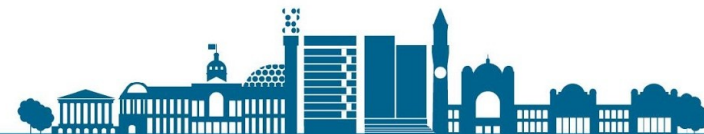
- ◆ Occurs when multiple processes or threads **read** and **write** shared data items
- ◆ The processes “**race**” to perform their read /write actions
- ◆ The final result depends on the order of execution
  - The “**loser**” of the race is the process that performs the last update and determines the final value of a shared data item






# Why do Race Conditions Occur?

- ◆ **Order of Execution:** Whenever the state of a shared resource depends on the precise execution order of the processes
- ◆ **Scheduling:** Context switches at arbitrary times during execution
- ◆ **Outdated Information:** Processes / Threads operate with **Stale** or **Dirty** copies of memory values in registers / local variables
  - Other processes may already have changed the original value in the shared memory location
  - **How can we avoid race conditions?**

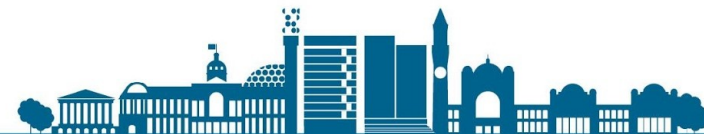


# Critical Section

- ◆ Part of the program code that accesses a shared resource



- ◆ In order to avoid race conditions, we have to control the concurrent execution of the **critical sections**
  - Strict Serialization i.e. **Mutual Exclusion**



# Critical Section

```
process ()  
{  
    entry_protocol()  
    critical_section()  
    exit_protocol()  
}
```

## ◆ Entry Protocol

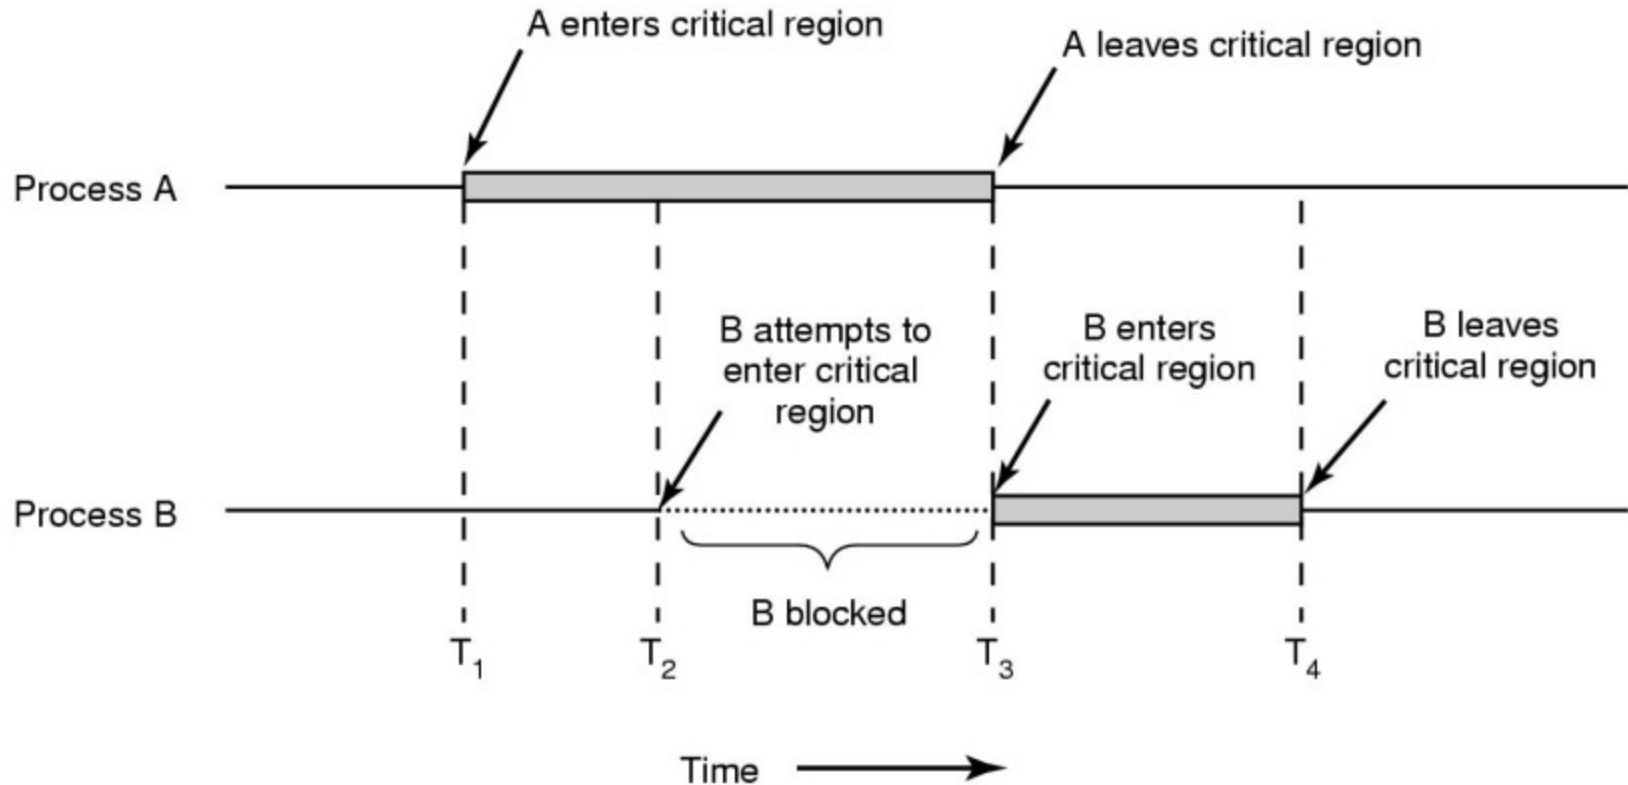
- Process requests permission for **entering a critical section**; May have to wait / suspended until entry is granted
- Process has to communicate that it **entered critical section**

## ◆ Exit Protocol


- Process communicates to other processes that it has **left the critical section**



## Critical Section / Mutual Exclusion



- ◆ **Important:** A process can finish the execution of its critical section, even if it is pre-empted or interrupted.





# Deadlock and Starvation

- ◆ Enforcing mutual exclusion creates two new problems
  - **Deadlocks** – Processes wait forever for each other to free resources
  - **Starvation** – A Process waits forever to be allowed to enter its critical section
- ◆ Implementing mutual exclusion has to account for these problems!



# Requirements for Solutions to the Critical Section Problem


- 1) **Serialization of Access:** Only one process at a time is allowed in the critical section for a resource
- 2) **Bounded Waiting (No Starvation):**
  - A process waiting to enter a critical section, must be guaranteed entry (within some defined limited waiting time)
  - Scheduling algorithm has to guarantee that process is eventually scheduled and can progress



# Requirements for Solutions to the Critical Section Problem (cont'd)

## 3) Progress (Liveness, No Deadlock)

- A process that halts in its noncritical section must do so without interfering with other processes currently waiting to enter their critical section
- Only processes currently waiting to enter their critical section are involved in the selection of the one process that may enter
- A process remains inside its critical section for a finite time only



# Solutions to the Critical Section Problem

## 1) Software Solutions

- Shared lock variables
- Busy Waiting (Polling / Spinning / Strict Alternation)

## 2) Hardware Solutions

- Disabling Interrupts
- Special H/W Instructions (like compare & swap)

## 3) Higher Level OS Constructs

- Semaphores, Monitors, Message Passing




# Lock Variables

- ◆ Critical sections must be protected by some form of a “lock”, where a lock is:
  - A shared data item
  - Processes have to “acquire” such a lock before entering a critical section
  - Processes have to “release” a lock when exiting critical section
- ◆ A lock is also called a **Mutex**.

mutual exclusion

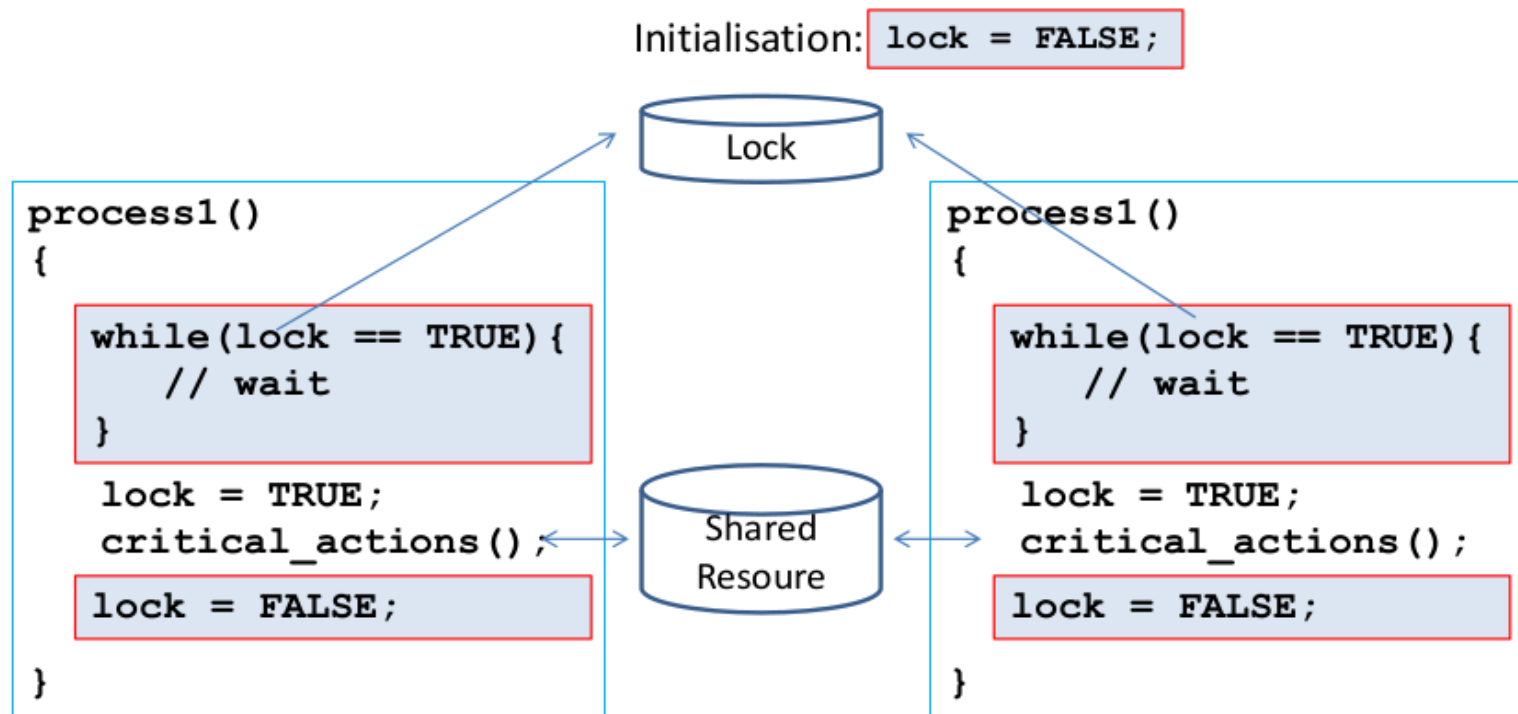
```
process ()  
{  
    acquire lock  
  
    critical_section() ;  
  
    release lock  
  
    remainder_section() ;  
}
```




# Shared Lock / Mutex

## ◆ Two State Lock (shared variable)

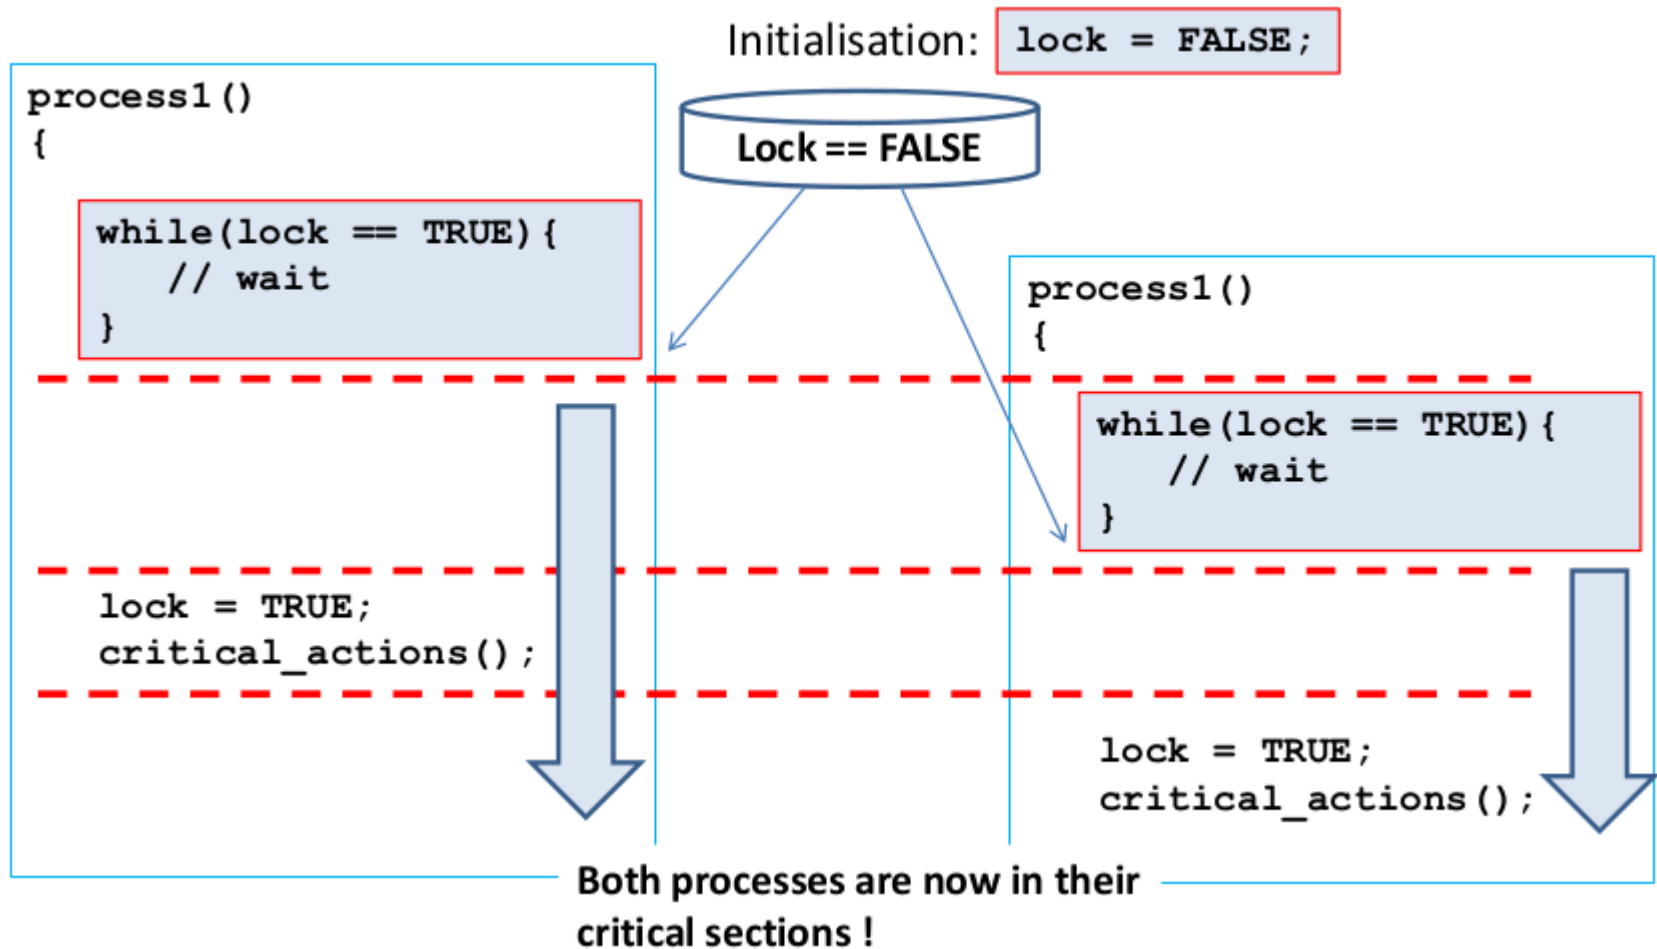
- **True** = Critical Section Locked, **False** = Critical Section Unlocked




**Problem:** As lock variable is itself a shared resource, race conditions can occur on it!



# Shared Lock - Problem

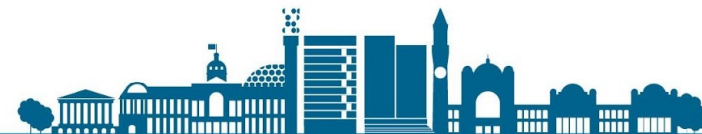


**Solution:** Use Atomic Instructions for Shared Lock!



## Busy Waiting (Polling, Spinning)

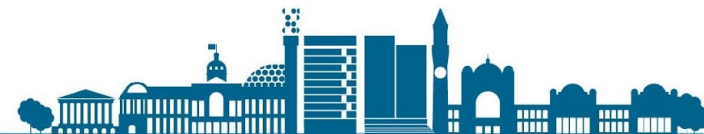
- ◆ A process **continuously evaluates** whether a lock has become available
- ◆ Lock is represented by a data item held in shared memory (IPC via shared memory)
- ◆ Process consumes CPU cycles without any progress
- ◆ A process busy waiting may prevent another process holding the lock from executing and completing its critical section and from releasing the lock
- ◆ **Spin locks** are used at kernel level (special HW instructions)



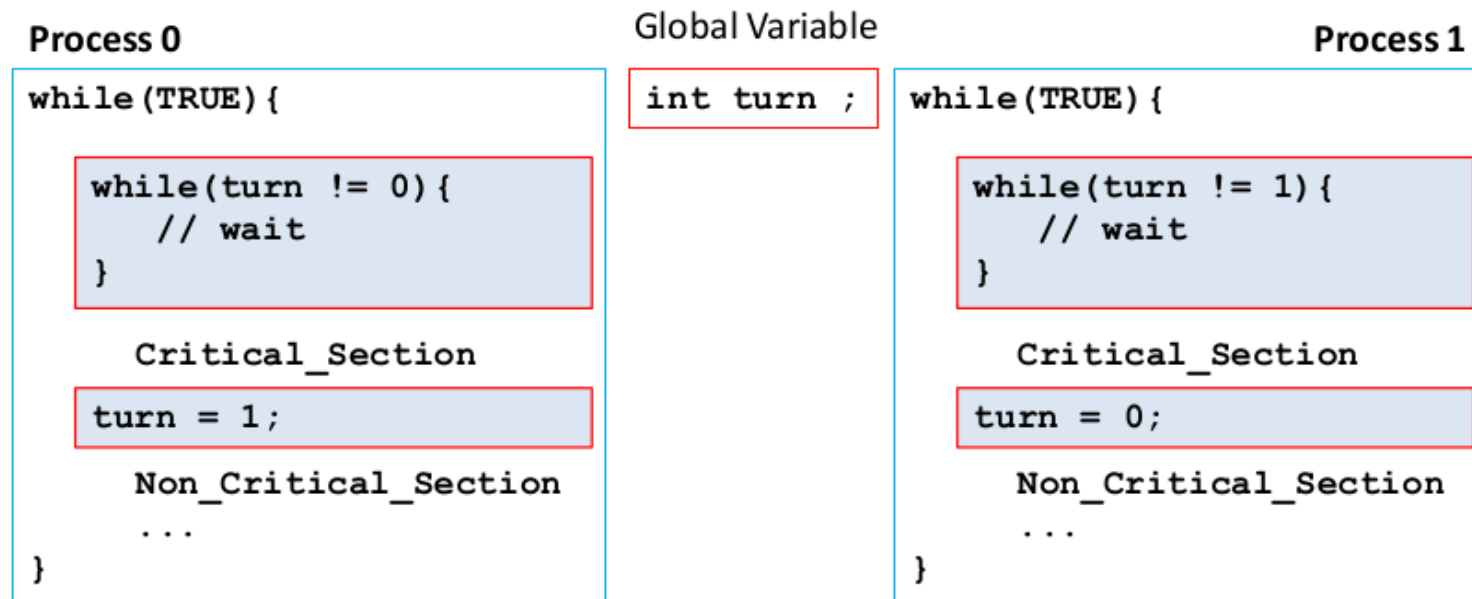


# Busy Waiting (Strict Alternation)


- ◆ Strict alternation between two processes – a **process waits** for its **turn**
- ◆ Use a “token” as shared variable – usually process ID
  - indicates which process is the next to enter critical section, set by previous process
- ◆ **Entry to critical section**
  - Process  $P_i$  busy-waits until  $\text{token} == i$  (its own process ID)
- ◆ **Exit from critical Section**
  - Process  $P_i$  sets token to next process ID



# Busy Waiting (Strict Alternation)



- ◆ Mutual exclusion is guaranteed!
- ◆ Liveness / Progression problem
  - Both process depend on a change of the “turn” variable
  - If one of the processes is held up in its non-critical section, it cannot do that and will block the other process



# Hardware Solutions – Disabling Interrupts

- ◆ In uniprocessor systems, **concurrent** processes cannot be overlapped, they can only be **interleaved**.
- ◆ **Disabling interrupts** guarantees mutual exclusion!
- ◆ However, the efficiency of execution could be noticeably degraded.
- ◆ This approach will not work in a multiprocessor architecture

```
while (true) {  
    /* disable interrupts */;  
    /* critical section */;  
    /* enable interrupts */;  
    /* remainder */;  
}
```



## Special H/W Instructions

- ◆ Applicable to **any number of processes** on either uniprocessor or multiprocessors sharing main memory
- ◆ **Simple** and **easy** to verify
- ◆ Can be used to **support multiple critical sections**
- ◆ **Busy-waiting** is employed
- ◆ **Starvation** and **Deadlocks** are possible!



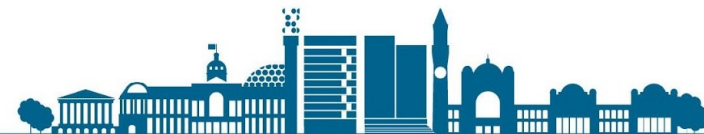
# Abstractions

- ◆ Usually, we build upon abstractions:
  - ◆ The system implements mechanisms correctly and efficiently and we use these:
    - ◆ Database systems
    - ◆ Python & Java
- ◆ There are other solutions to these problems that *may* be more efficient – in some circumstances



# Summary

- ◆ What is **concurrency** and what are **OS** and **Process concerns** w.r.t to concurrency.
- ◆ We have been introduced to the notion of **process synchronization**, its need and the **critical section problem**.
- ◆ We have seen **requirements** for **solutions to the CS problem** i.e. mutual exclusion, progress & bounded waiting time.




# Threads

- ◆ What is a thread?
- ◆ Parallelism vs. Concurrency
- ◆ Multicore and Multi-threading
- ◆ Multi-threading Models




# From Processes to Threads

- ◆ Process model assumed a process was an executing program with a **single thread** of control
- ◆ All modern OS's provide features enabling a process to contain **multiple threads** of control



- ◆ **Multi-threading** is the ability of an OS to support multiple, concurrent paths of execution *within* a single process.





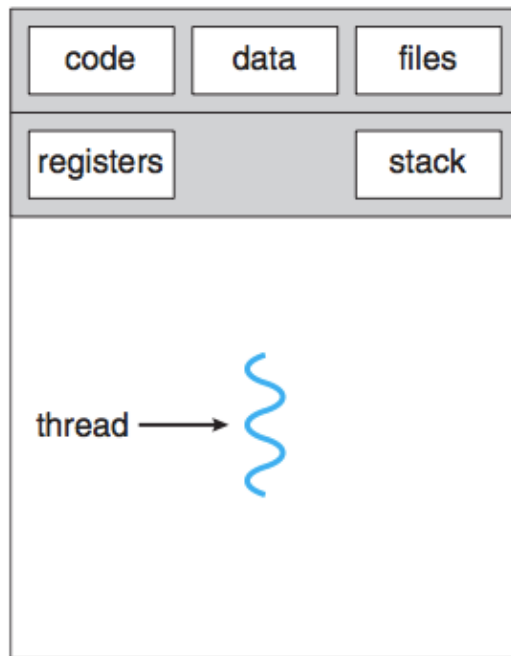
# What is a Thread?

- ◆ A thread is part of a process
- ◆ Contains:
  - Thread ID (TID)
  - Program Counter (PC)
  - Register Set
  - Stack
- ◆ All threads *in the same process*, share:
  - Code section
  - Data section:
    - Objects
    - Open files
    - Network connections
    - Etc.

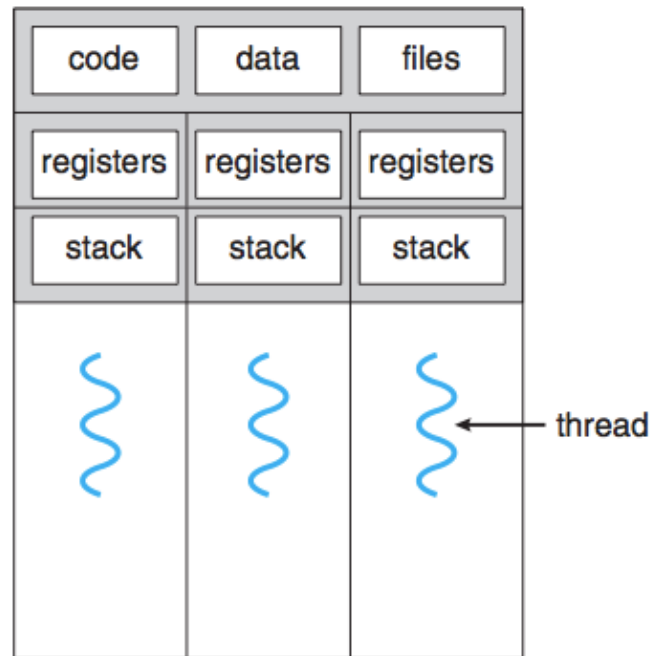


# What is a Thread?


- ◆ If a process has multiple threads of control, it can perform more than one task at a time!



single-threaded process



multithreaded process



# Why Threads?

Most applications *are* multi-threaded

## ◆ Application:

- ◆ Process with several threads of control

- ◆ Web browser:

  - ◆ Threads for each tab


  - ◆ Threads for each extension/helper

  - ◆ Utility threads

- ◆ Program: *main* thread + service threads

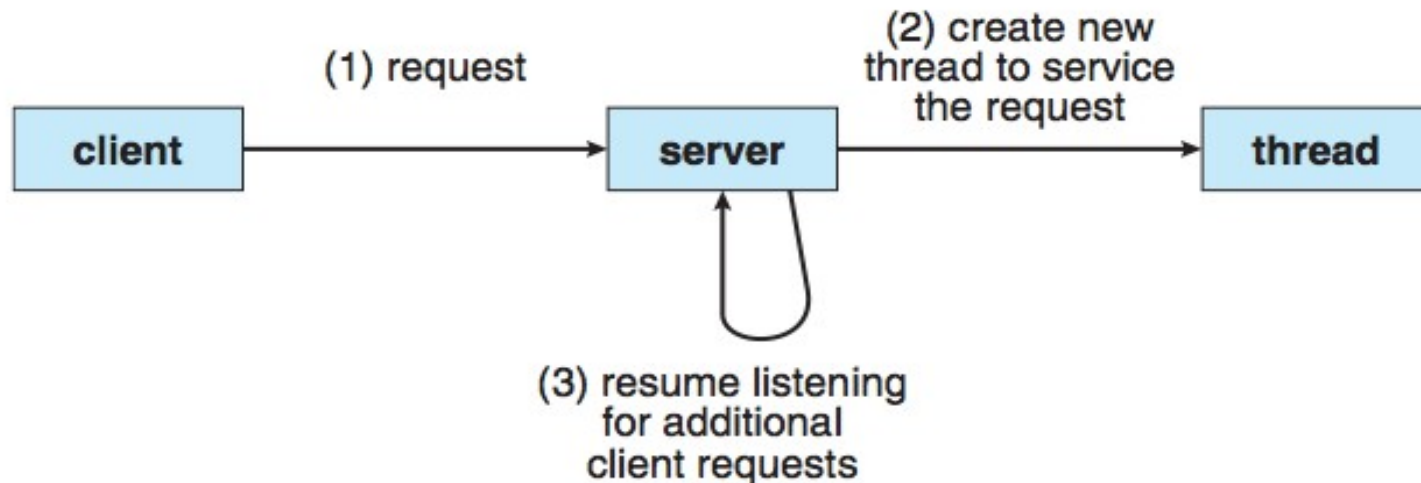
  - ◆ + GUI threads

- ◆ Server Processes



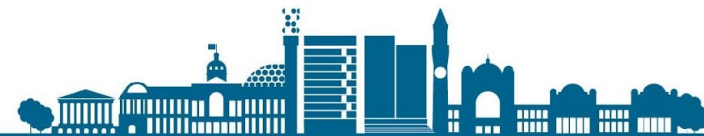
# Multi-threaded Server Architecture

- ◆ Threads have a critical role in systems which provide services:
  - ◆ Servers: web servers, database servers ....
- ◆ When a server receives a message, that request is (usually) serviced using a separate thread:



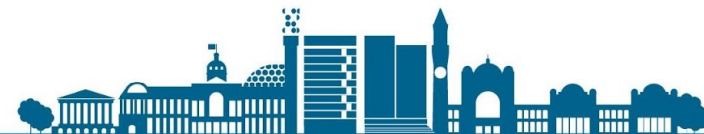

# Benefits of Threads

- ◆ **Responsiveness:** user interaction in a GUI can be responded to by a separate thread from one running a computationally intense algorithm.
- ◆ **Resource Sharing:** Variables & objects can be shared. This is especially important for things like network or DB connections.
- ◆ **Economy:** “Light-Weight Processes” due to smaller memory footprint than using multiple processes & efficiency of switching between threads.
- ◆ **Scalability:** Easier to make use of parallel processing in a multi-threaded processor
- ◆ **Reduce programming complexity:** Problems often are better broken down into independent parallel tasks, rather than having to build complex mechanisms to interleave their execution.




# Multicore Programming

- ◆ Multiple processing elements on a single chip
- ◆ Multi-threaded programming:
  - ◆ Allows efficient use of the multiple processing elements and improved parallelism

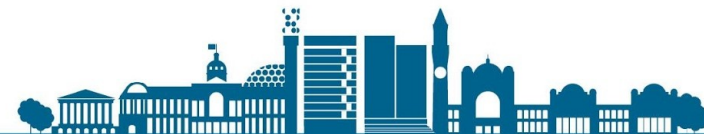


# Concurrency

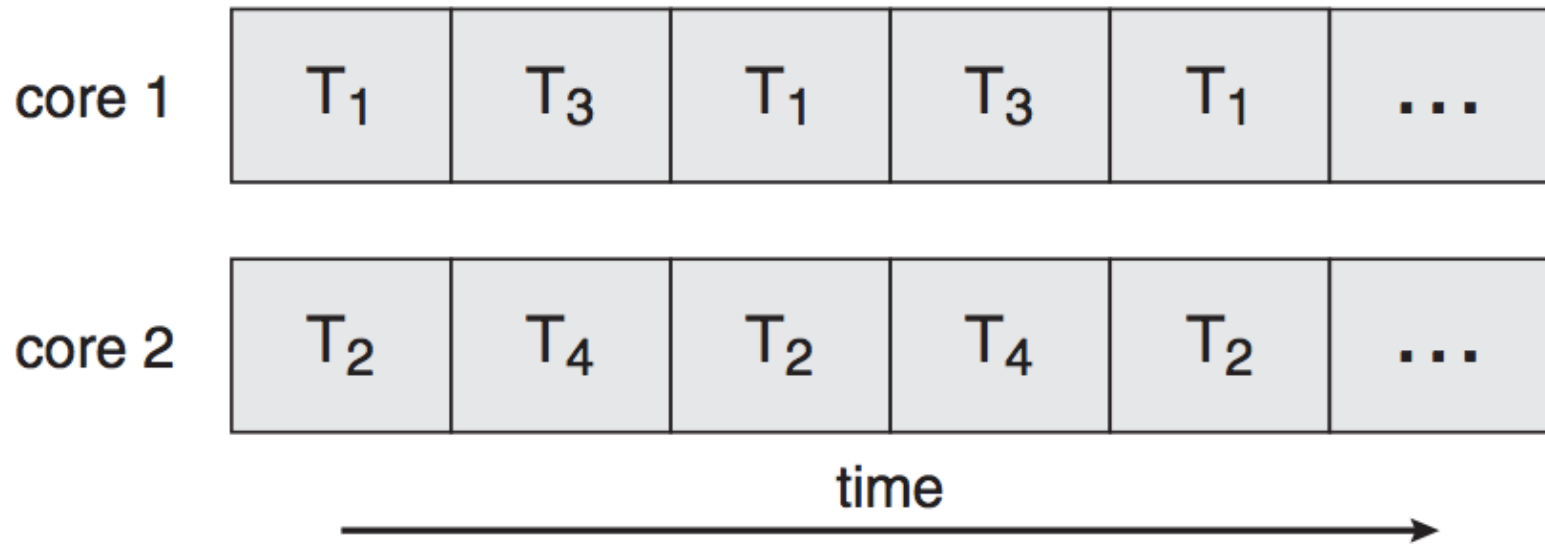
- ◆ On a single Processing element, **concurrency** means that the execution of threads will be interleaved over time.



- ◆ Only one thread at a time is actually running



# Parallelism






# Parallelism vs Concurrency

- ◆ A system is **parallel** if it can perform more than one task simultaneously
- ◆ A **concurrent** system supports more than one task by allowing all the tasks to make progress.
- ◆ It is possible to have concurrency without parallelism



This is  
incredible!  
I can do so  
much more!  
There must be a  
catch ...




# Multicore systems and Multithreading

- ◆ Performance of software on Multicore systems
- ◆ Amdahl's law states that:

$$\text{Speedup} = \frac{\text{time to execute program on a single processor}}{\text{time to execute program on } N \text{ parallel processors}} = \frac{1}{(1 - f) + \frac{f}{N}}$$

- ◆ **(1-f)** : Inherently serial code
- ◆ **f** : parallelizable code with no scheduling overhead



# Multicore and Multithreading (cont'd)

- ◆ Amdahl's law makes the multicore organizations **look attractive!**
- ◆ But even a small amount of **serial code** has noticeable impact on the overall performance


## **Example:**

10% serial, 90% parallel, 8 CPUs → ~4.7x speedup

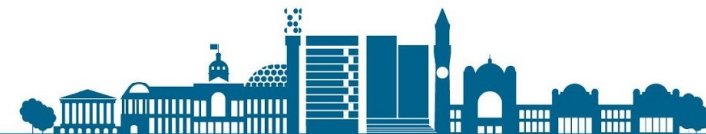
- ◆ Also: There are overheads to parallelism:
  - ◆ Memory access
  - ◆ Cache load
  - ◆ Etc.




## Multicore and Multi-threading (cont'd)




(a) Speedup with 0%, 2%, 5%, and 10% sequential portions



## Multicore and Multi-threading (cont'd)



(b) Speedup with overheads




# Parallel Programming Challenges

- ◆ Pressure placed on system designers and application programmers to **better use multiple cores**
- ◆ System designers must write scheduling algorithms that use multiple processing core to allow **parallel execution**
- ◆ Challenge to **modify existing programs**
- ◆ Challenge to **design new programs** that are multithreaded



# Parallel Programming Challenges (cont'd)

- ◆ **Identifying tasks:** which areas of an application can be divided into separate, concurrent (and ideally independent) tasks?
- ◆ **Balance:** how do we balance the tasks on the multiple cores to achieve maximum efficiency?
- ◆ **Data splitting:** how can data sets be split for processing in parallel?
- ◆ **Data dependency:** Certain tasks will rely on data from another. In this case, how do we synchronise the access to this data?
- ◆ **Testing and debugging:** Due to complexities of concurrent and parallel execution (multiple pathways etc.), how do we debug such an application?



# Multithreading Models


- ◆ Threads may either be at the **user** or **kernel** thread level
- ◆ User threads managed above kernel
- ◆ Kernel threads are managed directly by the OS
- ◆ Three types of relationships between user and kernel threads are possible.
  - Many-to-One
  - One-to-One
  - Many-to-Many






# Many-to-One Model

- ◆ Maps **many user-level** threads to **one kernel** thread
- ◆ Thread management done by the **thread library**, in user space. Therefore, efficient!
- ◆ But, entire process will block if a thread makes a blocking system call
- ◆ Only one thread can access kernel at a time - lacks parallelism on multicore systems
- ◆ Not many systems use this due to inability to take advantage of multiple cores



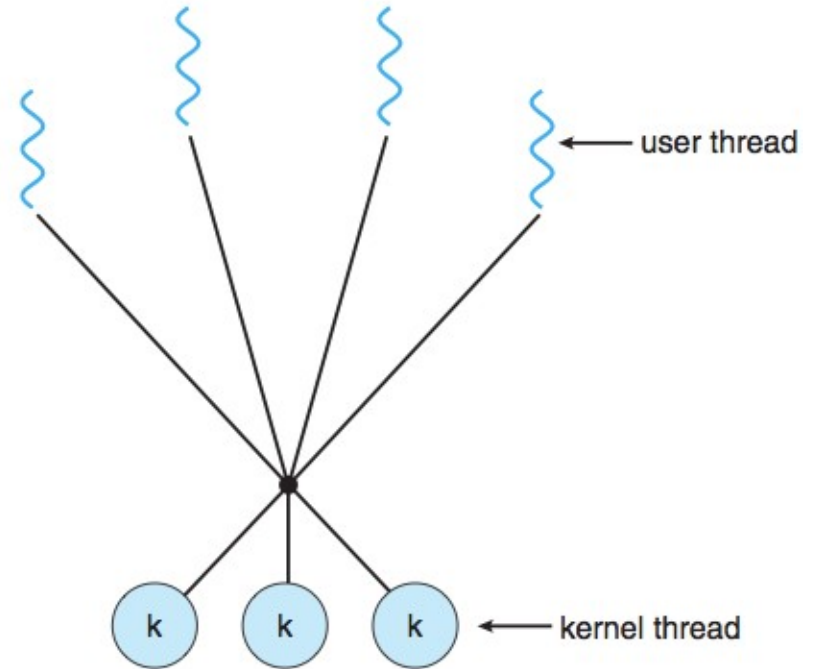
# One-to-One Model

- ◆ Maps **each user thread** to a **kernel thread**
- ◆ Overcomes blocking issue - so more concurrent
- ◆ Allows for parallelism
- ◆ But, for each user thread, there has to be a kernel thread



# Many-to-Many Model

- ◆ A software developer can create **as many user threads** as necessary i.e. a Thread Pool.
- ◆ Corresponding kernel threads can **run in parallel** on a multiprocessor
- ◆ If a thread blocks, kernel can schedule another for execution



# How can we use Threads?

## ◆ Through use of a **thread library**

- POSIX Pthreads (user/kernel level)
- Windows (kernel)
- Java Threads (on top of Windows/Pthreads)
- Python Threads
  - (See <https://realpython.com/intro-to-python-threading/#what-is-a-thread>)

## ◆ API for **creating** and **managing** threads



# Inter-thread Communication

- Threads communicate by sharing data:
  - They are part of the same process
  - **Efficient!**
- We now have the same problems that we discussed earlier!
- Solution?
  - **Use tools (within your programming language or package) to manage thread synchronisation**



## An Example

```
class Counter {  
    private int c = 0;  
  
    public void increment() {  
        c++;  
    }  
  
    public void decrement() {  
        c--;  
    }  
  
    public int value() {  
        return c;  
    }  
}
```



# Why is there a problem?

- ◆ Functions seem simple and atomic
- ◆ But, to increment, there are three steps
  - Retrieve the current value of integer
  - Add 1 to this value
  - Store the incremented value back to the variable



# Thread Interference

- ◆ If multiple threads all reference the same object:
  - ◆ **thread interference** can occur
- ◆ Happens when two functions, operating on the same data, **interleave**
- ◆ Unpredictable
- ◆ Difficult to debug





# Managing Threads and Thread interaction

- ◆ These topics will be covered next term:
  - ◆ Python thread management
  - ◆ Data base transactions
- ◆ These both provide effective, efficient and abstract ways to manage concurrent access to shared data



# Summary

In this part, we have introduced:

- ◆ The concept of threads and multithreading.
- ◆ The differences between parallelism and concurrency
- ◆ The Amdahl's law for performance on multicore
- ◆ Multithreading Models
- ◆ Threading Issues including thread interference & memory consistency



## References / Links

- ◆ Chapter # 4: Threads, Operating System Concepts (9th edition) by Silberschatz, Galvin & Gagne
- ◆ Chapter # 4: Threads, Operating System Internals and Design Principles (7th edition) by William Stallings
- ◆ <https://realpython.com/intro-to-python-threading/#what-is-a-thread>
- ◆ <https://docs.oracle.com/javase/tutorial/essential/concurrency/locksyntax.html>
- ◆ <http://tutorials.jenkov.com/java-concurrency/creating-and-starting-threads.html>

