

Introduction to Computer Systems

More on Numbers: Representing Real Numbers



Lecture Objectives

- ◆ To introduce how **real numbers** are **represented** in computer systems.
- ◆ To explain some of the limitations of floating point representations



Slide #2 of 25

Lecture Outline

More on Numbers:

- ◆ Binary Arithmetic using 2's Complement
- ◆ Fixed Point Decimal to Binary
- ◆ Fixed Point Binary to Decimal
- ◆ Fixed Point Arithmetic
- ◆ Floating Point Numbers
- ◆ Numerical Precision



Slide #3 of 25

Binary Arithmetic – Observations

- ◆ Add a "sign" bit – assume on the left
- ◆ Then +5 = 0101 and -5 = 1101
- ◆ We know that +5 + -5 = 0, so this should hold in binary!
- ◆ But:

$$\begin{array}{r} 0101 \\ 1101 \\ \hline 10010 \end{array}$$

We need to Fix It!

- ◆ We need to be a bit smarter about this ...



Slide #4 of 25

Binary Arithmetic – 2's Complement

- ◆ In **2's Complement** the convention is:
 - ◆ All positive numbers start with 0
 - ◆ All negative numbers start with 1
 - ◆ Negation is achieved by:
 - ◆ Flipping all the bits
 - ◆ Adding 1 to the least significant (right-most)
 - ◆ Lets see the same example again:
 - ◆ +5 = 0101 and -5 = 1011 (in 2C notation)

$$\begin{array}{r} 0101 \\ 1011 \\ \hline 10000 \end{array}$$

How about Binary
Arithmetic? Subtraction
etc.

This carry will be
discarded



Slide #5 of 25

How Real Numbers are Represented?

How do we represent real numbers in computer systems?

Two of the common ways to achieve this are:

- ◆ **Fixed Point:** binary point is fixed e.g.
1101101.0001001
- ◆ **Floating Point:** binary point floats to the right of the most significant 1 and an exponent is used e.g.
1.1011010001001 x 2⁶
- ◆ IEEE 754 - https://en.wikipedia.org/wiki/IEEE_754



Slide #6 of 25

Fixed Point Decimal-to-Binary

- ◆ Integer part convert as before (repeated division by 2)
- ◆ Non-integer part follows opposite process
- ◆ Repeated multiplication by 2, keeping integer part:

$0.537 \times 2 = 1.074$
 $0.074 \times 2 = 0.148$
 $0.148 \times 2 = 0.296$
 $0.296 \times 2 = 0.592$
 $0.592 \times 2 = 1.184$
 $0.184 \times 2 = 0.368$
 So $0.537_{10} = 0.100010_2$

- ◆ 0.625?
- ◆ 0.512?



Fixed Point Binary-to-Decimal

- ◆ Allocate subset of bits to integer part, and the remainder to the non-integer part.
- ◆ For example, 4+4 bits:

$$\begin{aligned}
 & 1101.0101 \\
 & 1x2^3 + 1x2^2 + 0x2^1 + 1x2^0 + 0x2^{-1} + 1x2^{-2} + 0x2^{-3} + 1x2^{-4} \\
 & = 8 + 4 + 0 + 1 + 0.0 + 0.25 + 0.0 + 0.0625 \\
 & = 13.3125
 \end{aligned}$$

- ◆ 101.110?
- ◆ 010.001?



Fixed Point Arithmetic

2's Complement

11101.011

whole numbers +1

- ◆ Everything is the same as for whole numbers
- ◆ Example: $01001.010 - 00010.100$
- ◆ Take 2C and add:

$$\begin{array}{r} 01001010 \\ + 11101100 \\ \hline (1)00110110 \end{array}$$

- ◆ 00110.101 – 10110.010?



Decimal Fractions (Base 10)

$$\begin{aligned} & 42 \cdot 625 \\ & \quad 10 \quad 1 \quad \frac{1}{10} \quad \frac{1}{100} \quad \frac{1}{1000} \\ & 4 \times 10 + 2 \times 1 + \frac{6}{10} + \frac{2}{100} + \frac{5}{1000} \\ & = 42625 \times 10^{-3} \\ & = 42 \frac{625}{1000} = 42 \frac{25}{40} = 42 \frac{5}{8} \end{aligned}$$



Infinite Decimal Fractions (Base 10)

e.g. $\frac{1}{3} = \underline{0.3333} \mid \underline{3333} \dots$

If only fixed number of decimal places allowed, the rest is lost

e.g. four places:

0.33333333 ...
rounding error



Binary Fractions (Base 2)

Similar. e.g.

Similar. e.g.

$$42\frac{5}{8} = 101010 \cdot 101 = 101010101 \times 2^{-3}$$

32 16 8 4 2 1 $\frac{1}{2}$ $\frac{1}{4}$ $\frac{1}{8}$

binary

$$\frac{1}{10} = 0.000110011001100\dots$$

decimal

infinite binary fraction



Floating Point Representation (for fractions)

General principle

- like "scientific notation", but in binary

Numbers represented as: $\pm m \cdot 2^e$

Sign (\pm)	Mantissa (m)	Exponent (e)
1 bit 0 for + 1 for -	Actual significant digits	'Signed' binary 'Shows' where binary point goes



Slide #13 of 25

Choice of Representations

$$\begin{aligned}
 42 \frac{5}{8} &= 101010.101 = 1010101 \times 2^{-3} \\
 &= 1010101010 \times 2^{-4} \\
 &= 10101010100 \times 2^{-5} \\
 &\approx \dots \\
 \text{or} \quad &= 101010.101 \times 2^0 \\
 &= 10101.0101 \times 2^1 \\
 &= \dots \approx 1.01010101 \times 2^5
 \end{aligned}$$



Slide #14 of 25

Floating Point Representation in Java

S	Offset e	Mantissa m
---	----------	------------

Sign: 0, 1 for +, -
Exponent: Biased binary
Mantissa: Leading mantissa bit (1.) left out

- ◆ take bits of m
- ◆ put "1." at start, so normalized
- ◆ move binary point right e places
- ◆ (or left for negative e)
- ◆ note - e is stored with an "offset" added to it



Slide #15 of 25

Java Types for Floating Point

Type	No of Bits				Bytes
	Sign	Mantissa	Exponent	Total	
float	1	23	8	32	4
double	1	52	11	64	8

1 extra for hidden bit

52 bit mantissa : $2^{53} \approx 8 \times 10^{15}$

We get 15 significant decimal digits in double data type.



Slide #16 of 25

Java Types for Floating Point

Offset: $2^{(n-1)} - 1$

$2^{(8-1)} - 1 = 127$

$$42 \frac{5}{8} = 101010.101 = 1.01010101 \times 2^5$$

0	1000 0100	010 1010 1000 0000 0000 0000
---	-----------	------------------------------

Sign: 0 for +
5 + offset 127 = 132
e + 127
Mantissa: Without the leading 1.

$$\begin{aligned}
 \frac{1}{10} &= 0.000110011001100\dots \\
 &= 1.10011001100\dots \times 2^{-4}
 \end{aligned}$$

0	0111 1011	100 1100 1100 1100 1100 1101
---	-----------	------------------------------

Sign: 0 for +
-4 + offset 127 = 123
e + 127
Mantissa: Without the leading 1.

Rounded Up
Rounding Error!



Slide #17 of 25

Money as Floating Point?

Floating Point value for amount in pounds with pence as fraction?

Not a good practice!

Pence need infinite binary fractions
e.g. 10p is 0.0001100110011001100...
so we get rounding errors

Always use int or long (or BigInteger) for money.



Slide #18 of 25

Factorial with Double

```
/**
 * Calculate factorial.
 * requires: 0 <= n
 * @param n number whose factorial is to be calculated
 * @return factorial of n
 */
public static double dfact(int n){
    double a = 1;
    for (int i = 1; i <= n; i++){
        a = a * i;
    }
    return a;
}
```

n, n!
165, 5.423910666131586E295
166, 9.003691705778433E297
167, 1.5036165148649983E300
168, 2.526075744973197E302
169, 4.2690680090047027E304
170, 7.257415615307994E306
171, Infinity
172, Infinity
173, Infinity
174, Infinity

Slide #19 of 25

Accuracy Issues – Even in Double Floating Point

n = 170
 n! = 7.257415615307994E306
 = 7257415615307994E291

This digit is wrong!

291 more digits after it

Windows calculator says:

7.2574156153079989673967282111293e+306

- ◆ Floating point arithmetic loses accuracy in least significant digits.
- ◆ Most significant, and overall size, are OK.

Slide #20 of 25

Why is 171! too big?

$170! \approx 7.2 \times 10^{306}$
 $171! \approx 171 \times 7.2 \times 10^{306}$
 $\approx 1231 \times 10^{306}$
 $\approx 10^{309} = (10^3)^{103} \approx 2^{1030}$
 Needs binary exponent $\approx 1030 > 1024 = 2^{10}$

Exponent too big to fit in 11 bits (signed) for Java double

Slide #21 of 25

Floating Point Overflow

In Java:

If result is too big for datatype,

it's a special value **POSITIVE_INFINITY**

More precisely:
 Float.POSITIVE_INFINITY
 or
 Double.POSITIVE_INFINITY

Other special values:

If too big but negative: **NEGATIVE_INFINITY**

Indistinguishable from 0 but known to be negative: -0.0

Impossible number (e.g. Sqrt(-1)) **NaN** "Not a Number"

These special values allow you to check for overflow in a program

- unlike the case for integer arithmetic

Slide #22 of 25

Summary – Java Floating Point

- ◆ It normalizes where possible
- ◆ **Unnormalised for smallest numbers (offset e = 0)**
- ◆ Special representations for NaN etc.
- ◆ Details in API for:
 - java.lang.Float.intBitsToFloat
 - java.lang.Double.longBitsToDouble

Slide #23 of 25

Numerical Precision

- ◆ Fixed point is convenient and intuitive but has two problems

1) Numerical precision

- Only values that are an integer multiple of the smallest power of two can be represented exactly

2) Numerical range

- Increased precision of non-integer part is at the expense of numerical range

- ◆ Floating point representation effectively addresses these issues.

Slide #24 of 25

Summary – Numbers

- ◆ Representing numbers in the computer
 - Whole numbers in binary and hexadecimal notation
 - Positive real numbers in fixed-point binary
- ◆ Binary arithmetic is like decimal arithmetic
 - Our lack of binary practice makes it hard!
- ◆ Negative numbers are tricky things
 - But we can use a few of our own tricks – 2C
- ◆ Floating point is an alternative, but is very unnatural for us!

