



UNIVERSITY OF
BIRMINGHAM

Computer Systems Algorithm Design & Analysis: Efficiency & Complexity




Motivation

We want to build computer systems that work

- And that work well!

To do this we need to think about several factors:

- They must give the correct answer
 - Is this ALWAYS true?
- They must be reliable, maintainable, quick to produce, cheap ...
 - Again, is this always true?
- They must be usable – by the user of the system
- They must be efficient



Overview

Algorithm Design & Analysis

- ◆ Motivation
- ◆ Efficiency and Space complexity
- ◆ Efficiency: Time complexity
 - ◆ Of algorithms
 - ◆ Of problems
- ◆ Conclusions and lessons



Efficient?

We want our software to:

- Run as quickly as possible
- Respond as quickly as possible
- Use as little memory as possible
- Use as little network bandwidth as possible
- Use as little power as possible
-



Efficient?

So, we are actually trying to optimise many things:

- The balance will depend on our problem

One serious issue we need to consider is:

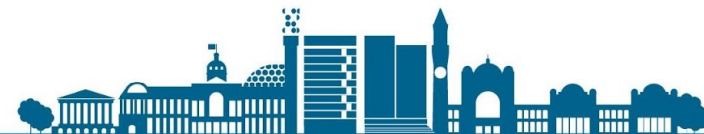
- Our solution may work very well for 'small' problems
- But, if the problem gets more complex how does its performance change?
 - Does it get worse?
 - Does it get MUCH worse?
 - Does it get MUCH, MUCH worse?



What are we interested in?

Actually, there are several things to consider:

- Is our implementation of our algorithm or data representation efficient:
 - Could we improve our code to make it run faster?
 - This might speed things up a bit
 - or a lot ... maybe 10 or 100 times
 - Could we improve our data representation so it uses less memory?
 - Again we might improve things dramatically
- But what if our *algorithm* is inherently going to perform worse as the amount of data increases?
- What if our *problem* is inherently hard and gets much harder as the size of the problem grows?



So?

We need to understand how our problem and/or algorithm changes in complexity as it gets larger.

Usually, we focus more on time complexity – how much computation is involved in solving a problem:


- For instance, if we *double* the size of our problem (e.g. twice as much data) then what happens to the computation time:
 - Does it stay the same?
 - Does it get a bit harder? Maybe +30%
 - Does it get a lot harder? Maybe +100%
 - Does it get even worse? Maybe +400% or +800% or worse?
- Usually, we are interested in an almost *qualitative* measure



Space complexity

Usually, space (or bandwidth) complexity is not considered in the same way as computational complexity:

- Memory and bandwidth are limited by practical constraints
- We can add memory (and bandwidth or processing) to cope with a bigger problem
- We can constrain our problem to the practical constraints
 - We restrict the resolution of videos to what is practical
 - We restrict the resolution of 3D models to what is practical
 - We can use other techniques to reduce the memory demands (e.g. compression or streaming)




An Example: video

Typical HD video:

- 1920x1080 pixels (4 bytes per pixel):
 - ~2MPixels, ~8Mbytes
- 25 frames per second
- 200 Mbytes/sec (~1.6 Gbps connection)!

Solution:

- Compression (up to 200:1)
 - Using empirical knowledge of human vision & video streams
 - Lossy compression (cf. lossless compression)
- Reduce frame rate or resolution
- Constrained by practical constraints



How does the problem change with resolution?

If we double the resolution of a 2D image

- We x4 the number of pixels (n^2)
- If we have a 3D image (e.g. from an MRI scanner)
 - If we double the resolution then: x8 (n^3)

The same problem will occur with 2D models and will be even worse as the number of dimensions increases:

- n^2 , n^3



Time complexity

To develop a *basic* understanding
of algorithm *design* and *efficiency* in terms of
worst case time complexity.




Outline / Topics

- ◆ Algorithm Design and Analysis
- ◆ Efficiency
- ◆ Time Complexity
- ◆ Big-O Notation
- ◆ Examples
- ◆ Lessons




What is an Algorithm?

In mathematics, computer science, and related subjects, an algorithm is an effective method for solving a problem expressed as a finite sequence of instructions.



- Abu Abdullah Muhammad bin Musa Al-Khwarizmi
 - a Muslim mathematician
 - invented *algorithms* to solve quadratic equations
 - the word *algorithm* derived from his Latin name *Algorithmi*
- An even earlier *algorithm* is the *sieve of Eratosthenes*



Algorithm: question?

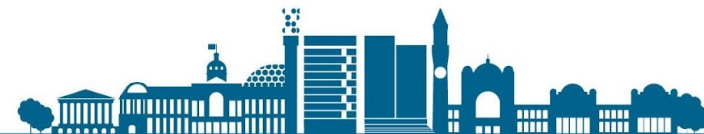
Which one is an algorithm?

- ◆ A recipe for making tomato soup?
- ◆ A procedure to sort 1000 numbers into ascending numeric order?
- ◆ A procedure to recognize a particular face in a crowd?
- ◆ A method to order objects according to beauty?



Algorithm Design and Analysis

- ◆ Multiple algorithms often exist for the same task:
 - All of them give correct results
 - How do we select the best one?
- ◆ Many possible (and often conflicting) criteria:
 - Efficiency
 - ?
 - simplicity, clarity
 - elegance, proofs of correctness
- ◆ We need to ask
 - is my algorithm correct?
 - does my algorithm always terminate?
 - does an algorithm even exist?



Algorithm Efficiency

◆ Resource usage of an algorithm

- typically: **time** (runtime) and **space** (computer memory)
- also: network usage, hardware requirements, ...
- consider trade-offs between resources

◆ How do we measure the run-time of an algorithm?

- benchmarking on *representative* set of inputs: **empirical analysis**
- analyse the **(time) complexity**



Algorithm Efficiency – Empirical Analysis


◆ Idea: Implement the algorithm (a program) and time it

Question: How can we time a program?

- Manual: Using stopwatch
- Automatic: Using some timer function

Run the program for various input sizes and measure run time.

data size	time (ms)
250	5
500	8
1000	10
2000	15
4000	30
8000	50
16000	75



Algorithm Efficiency – Time Complexity

◆ Time complexity:


- the number of **operations** that an **algorithm** requires to execute, **in terms of** the **size** of the input or problem

Note:

- *algorithm*, not implementation
 - so: *pseudocode*; no fixed programming language, computer architecture
- "in terms of" – complexity defined as a function **T(n)**

◆ Questions:

- what do we mean by *operations*?
- what do we mean by *size*?
- we usually focus on **worst-case**, not **average-case**, analysis



Example # 1

- ◆ Look up a value v in an array x of integers

1	4	17	3	90	79	4	6	81
---	---	----	---	----	----	---	---	----

- ◆ Algorithm: Linear Search

- inputs: array x of size n , integer v
- return: index of first occurrence of v in x , or -1 if none

- ◆ $T(n) = n$

```
for i=0...n-1:
    if x[i] == v:
        return i
return -1
```

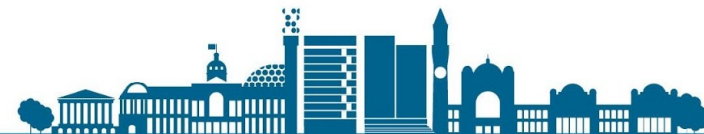


Example # 2

- ◆ Matrix-Vector Multiplication: $\mathbf{x} = \mathbf{A} \mathbf{b}$
 - $n \times n$ matrix \mathbf{A} , vector of \mathbf{b} size n
- ◆ Algorithm
 - inputs: matrix \mathbf{A} , vector \mathbf{b}
 - result stored in vector \mathbf{x} (initially all 0)

```
for i=0...n-1:  
    for j=0...n-1:  
         $\mathbf{x}[i] = \mathbf{x}[i] + \mathbf{A}[i][j] * \mathbf{b}[j]$ 
```

◆ $T(n) = 2n^2$




Quadratic Growth

Consider the two functions

$$f(n) = n^2$$


$$g(n) = n^2 - 3n + 2$$

Around $n = 3$, they look very different




Quadratic Growth

Around $n = 10$, they look similar




Quadratic Growth

Around $n = 100$, they are almost the same



Quadratic Growth

Around $n = 1000$, the difference is indistinguishable!



Quadratic Growth



The absolute difference is large, such that,

$$f(1000) = 1000000$$

$$g(1000) = 997002$$

but the relative difference is very small,

$$\left| \frac{f(1000) - g(1000)}{f(1000)} \right| = 0.002998 < 0.3\%$$

and this difference goes to zero as $n \rightarrow \infty$



Big-O Notation

◆ Usually we don't need exact complexity $T(n)$

- it suffices to know the **complexity class**
- we ignore constant factors/overheads, lower orders
- focus on performance for **large n** ("asymptotic")

◆ Big-O notation

- for example: $O(n^2)$ "O of n squared" "on the order of n^2 "

◆ Examples

- $T(n) = n \implies$ complexity class = $O(n)$
- $T(n) = n+2 \implies$ complexity class = $O(n)$
- $T(n) = 2n^2 \implies$ complexity class = $O(n^2)$



Big-O Notation

◆ More Examples

- $T(n) = 10n^3 + 1 \implies$ complexity class = $O(n^3)$
- $T(n) = 5(n+2) \implies$ complexity class = $O(n)$
- $T(n) = 1000 \implies$ complexity class = $O(1)$ 💬
- $T(n) = n^2 + n + 1 \implies$ complexity class = $O(n^2)$


◆ Determining the complexity class

- intuitively, it suffices to count the number of loops and the number of times they are executed



Big-O Notation: Common Complexity Classes


◆ Some common complexity classes:

- $O(1)$ = “Constant”
- $O(\log_2 n)$ = “Logarithmic” 
- $O(n)$ = “Linear”
- $O(n \log_2 n)$ = “Log Linear”
- $O(n^2)$ = “Quadratic”
- $O(n^3)$ = “Cubic”
- $O(2^n)$ = “Exponential”



◆ **Polynomial:** $O(n)$, $O(n^2)$, $O(n^3)$, ... – “tractable”


◆ **Exponential:** $O(2^n)$, $O(C^n)$, – “intractable”



Some concrete numbers ...

- ◆ How many **operations** needed for an algorithm?
 - with complexity $T(n) = f(n)$ and input size n

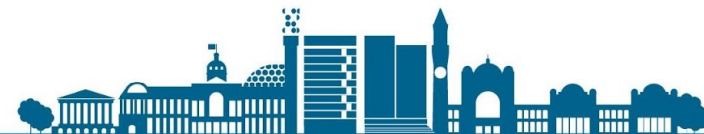
$f(n)$	$n = 4$	$n = 16$	$n = 256$	$n = 1024$	$n = 1048576$
1	1	1	1	1.00×10^0	1.00×10^0
$\log_2 \log_2 n$	1	2	3	3.32×10^0	4.32×10^0
$\log_2 n$	2	4	8	1.00×10^1	2.00×10^1
n	4	16	256	1.02×10^3	1.05×10^6
$n \log_2 n$	8	64	2.05×10^3	1.02×10^4	2.10×10^7
n^2	16	256	6.55×10^4	1.05×10^6	1.10×10^{12}
n^3	64	4.10×10^3	1.68×10^7	1.07×10^9	1.15×10^{18}
2^n	16	65536	1.16×10^{77}	1.80×10^{308}	6.74×10^{315652}




Some concrete numbers ...

- ◆ How much **time** needed for an algorithm?
 - assuming 1 million operations per second

$f(n)$	$n = 4$	$n = 16$	$n = 256$	$n = 1024$	$n = 1048576$
1	1 usec	1 usec	1 usec	1 usec	1 usec
$\log_2 \log_2 n$	1 usec	2 usec	3 usec	3.32 usec	4.32 usec
$\log_2 n$	2 usec	4 usec	8 usec	10 usec	20 usec
n	4 usec	16 usec	256 usec	1.02 msec	1.05 sec
$n \log_2 n$	8 usec	64 usec	2.05 msec	10.2 msec	21 sec
n^2	16 usec	256 usec	65.5 msec	1.05 sec	12.72 days
n^3	64 usec	4.1 msec	16.8 sec	17.83 min	36559 yrs
2^n	16 usec	65.5 msec	3.7×10^{63} yrs	5.7×10^{294} yrs	2.14×10^{315639} yrs



Graphically



Example # 1 (again)

- ◆ Look up a value **v** in an array **x**

1	4	17	3	90	79	4	6	81
---	---	----	---	----	----	---	---	----



Example # 1 (again)

- ◆ Look up a value **v** in a sorted array **x**

1	3	4	4	6	17	79	81	90
---	---	---	---	---	----	----	----	----



Example # 1 (again)

- ◆ Look up a value v in a sorted array x


1	3	4	4	6	17	79	81	90
---	---	---	---	---	----	----	----	----

- ◆ **Algorithm: Binary Search**

- inputs: sorted array x of size n , integer v
- return: index of first occurrence of v in x , or -1 if none

- ◆ **Complexity** = $O(\log_2 n)$

```
left = 0; right = n-1
while left < right:
    mid = (left+right)/2
    if x[mid] < v:
        left = mid+1
    else:
        right = mid
if x[left] == v:
    return left
else:
    return -1
```



Computing Complexity Classes

- ◆ Determine total algorithm complexity
 - from the complexities of its components

- 1) Sequential algorithm phases
- 2) Function / Method calls



Sequential Algorithm Phases

- ◆ Example: matrix-vector multiplication: $x = A b$
 - matrix-vector multiplication $x = A b$
 - initialise vector x (all 0), then add values to x

```
for i=0...n-1:  
    x[i] = 0  
for i=0...n-1:  
    for j=0...n-1:  
        x[i] = x[i] + A[i][j] * b[j]
```

- ◆ Complexity: $O(n) + O(n^2) = O(n^2)$
 - In general: “maximum” of complexities



Function/Method calls

- ◆ Example: n array look-ups

```
for i=0... $n$ -1:  
    binary_search( $x$ ,  $v_i$ )
```

- ◆ Complexity: $O(n) \times O(\log n) = O(n \log n)$
- ◆ In general: “multiplication” of complexities



Computing Complexity Classes

- ◆ Determine total algorithm complexity
 - from the complexities of its components

1) Sequential algorithm phases: “maximum”

- e.g. $O(n) + O(n^2) = O(n^2)$
- e.g. $O(n) + O(\log n) = O(n)$

2) Function / Method calls: “multiplication”

- e.g. $O(n) \times O(\log n) = O(n \log n)$
- e.g. $O(n^2) \times O(1) = O(n^2)$



Question # 1

Given the following code fragment, what is its Big-O running time?

```
1  test = 0
2  for i in range(n):
3      for j in range(n):
4          test = test + i * j
```

- a) $O(n)$
- b) $O(n^2)$
- c) $O(\log n)$
- d) $O(n^3)$



Question # 2

Given the following code fragment, what is its Big-O running time?

```
1  test = 0
2  for i in range(n):
3      test = test + 1
4
5  for j in range(n):
6      test = test - 1
```

- a) $O(n)$
- b) $O(n^2)$
- c) $O(\log n)$
- d) $O(n^3)$




Question # 3

Given the following code fragment, what is its Big-O running time?

```
1  i = n
2  while i > 0:
3      k = 2 + 2
4      i = i / 2
5  }
```

/ 0 0 0
5


- a) $O(n)$
- b) $O(n^2)$
- c) $O(\log n)$
- d) $O(n^3)$



Some Harder Problems

◆ Traveling Salesman Problem (TSP)

- given n cities and the distances between them, what is the **shortest possible route** that **visits each city exactly once** and then **returns to the first city**?



“A naive approach to solving TSP would be the brute-force solution, which means finding all possible routes given a certain number of nodes. This is a very expensive way to solve it, with a time complexity of $O(n!)$.”

<https://tonicanada.medium.com/introduction-to-the-travelling-salesman-problem-5ace44932cb5>

<https://www.youtube.com/watch?v=SC5CX8drAtU>




Some Harder Problems

◆ Boolean Satisfiability Problem (SAT)

- Given a Boolean formula $F(x_1, x_2, x_3, \dots, x_n)$
- Can F evaluate to 1 (true)?
- If yes, return values of x_i 's (satisfying the assignment) that make F true.

◆ In both cases (TSP & SAT):

- lots of practical applications (Operations Research, Optimization, Logistics, Model Checking, Software Verification, ... etc.)
- also important problems in theoretical computer science




Algorithms vs. Problems

◆ Algorithm complexity

- worst-case run-time of algorithm – "efficiency"
- actually an upper bound: algorithm A is in $O(f(n))$ if the worst-case run-time is at most $f(n)$
- usually use tightest (most informative) complexity

◆ Problem complexity

- complexity class = set of problems
- problem X is in complexity class $O(f(n))$ if there exists an algorithm to solve it in $O(f(n))$ – "difficulty"
- again: this is an upper bound (and uses the tightest possible)
- sometimes consider lower bounds:
- e.g. sorting: $O(n \log n)$



P and NP

◆ Complexity classes

- $O(1) \subseteq O(\log n) \subseteq O(n) \subseteq O(n \log n) \subseteq O(n^2) \subseteq O(2^n)$
- polynomial time (PTIME or P) – assumed to be "tractable"

◆ Another famous class: NP (Non-deterministic polynomial time)

- if we can "guess" a solution, it can be checked efficiently (efficiently = in polynomial time). e.g. Sudoku, Chess (more [at this link](#))
- In contrast, polynomial time (or P) problems are those where finding the answer is easy.


◆ NP-hard problems

- e.g. travelling salesman problem, SAT, ...
- only exponential time algorithms are known
- but some efficient heuristics exist in practice

◆ $P = NP?$

- ◆ Nobody knows ...
- ◆ [US\\$1,000,000 prize if you solve it!](#)

				3	7	6		
			6				9	
		8						4
	9							1
6								9
3							4	
7						8		
	1				9			
		2	5	4				



Summary

◆ Algorithm design and analysis:


■ Efficiency

◆ Time Complexity

■ (worst-case) number of operations an algorithm needs to execute, in terms of the size of the input or problem: $T(n)$

◆ Big-O notation

- complexity classes: $O(n)$, $O(n \log n)$, $O(2^n)$, ...
- focus on large values of n (i.e. asymptotic behaviour)
- ignore constants, lower factors
- count loops/iterations, decompose algorithm
- complexities for algorithms/problems



Lessons

- ◆ Data Representation/Algorithms are tightly linked
- ◆ Some algorithms are just better than others:
 - ◆ But sometimes the ‘best’ depends on the problem
- ◆ Some problems are inherently hard (in the complexity sense)
 - ◆ This doesn’t mean they cannot be solved
 - ◆ We may need to reformulate the problem
 - ◆ Constrain it
 - ◆ Relax conditions
 - ◆ E.g. instead of “find the best” we “find a very good”



Books / References

- ◆ Goldschlager and Lister: Computer Science: a Modern Introduction (Prentice Hall, 2nd edition 1988)
 - Chapter 3
- ◆ Aho & Ullman: Foundations of Computer Science (Freeman, 1992)
 - Chapter 3
 - <http://infolab.stanford.edu/~ullman/focs.html>

