



UNIVERSITY OF
BIRMINGHAM

Computer Systems

Algorithm Design & Analysis: Efficiency & Complexity




Motivation

We want to build computer systems that work

- And that work well!

To do this we need to think about several factors:

- They must give the correct answer
 - Is this ALWAYS true?
- They must be reliable, maintainable, quick to produce, cheap ...
 - Again, is this always true?
- They must be usable – by the user of the system
- They must be efficient



Overview

Algorithm Design & Analysis

- ◆ Motivation
- ◆ Efficiency and Space complexity
- ◆ Efficiency: Time complexity
 - ◆ Of algorithms
 - ◆ Of problems
- ◆ Conclusions and lessons



Efficient?

We want our software to:

- Run as quickly as possible
- Respond as quickly as possible
- Use as little memory as possible
- Use as little network bandwidth as possible
- Use as little power as possible
-



Efficient?

So, we are actually trying to optimise many things:

- The balance will depend on our problem

One serious issue we need to consider is:


- Our solution may work very well for 'small' problems
- But, if the problem gets more complex how does its performance change?
 - Does it get worse?
 - Does it get MUCH worse?
 - Does it get MUCH, MUCH worse?



What are we interested in?

Actually, there are several things to consider:

- Is our implementation of our algorithm or data representation efficient:
 - • Could we improve our code to make it run faster?
 - This might speed things up a bit
 - or a lot ... maybe 10 or 100 times
 - • Could we improve our data representation so it uses less memory?
 - Again we might improve things dramatically
- But what if our *algorithm* is inherently going to perform worse as the amount of data increases?
- What if our *problem* is inherently hard and gets much harder as the size of the problem grows?



So?

We need to understand how our problem and/or algorithm changes in complexity as it gets larger.


Usually, we focus more on time complexity – how much computation is involved in solving a problem:

- For instance, if we double the size of our problem (e.g. twice as much data) then what happens to the computation time:

- • Does it stay the same?
- • Does it get a bit harder? Maybe +30%
- Does it get a lot harder? Maybe +100%
- Does it get even worse? Maybe +400% or +800% or worse?

- Usually, we are interested in an almost qualitative measure






Space complexity

Usually, space (or bandwidth) complexity is not considered in the same way as computational complexity:

- Memory and bandwidth are limited by practical constraints
- We can add memory (and bandwidth or processing) to cope with a bigger problem
- We can constrain our problem to the practical constraints
 - We restrict the resolution of videos to what is practical
 - We restrict the resolution of 3D models to what is practical
 - We can use other techniques to reduce the memory demands (e.g. compression or streaming)




An Example: video

Typical HD video:

- • 1920x1080 pixels (4 bytes per pixel):
 - ~2MPixels, ~8Mbytes
- • 25 frames per second
 - 200 Mbytes/sec (~1.6 Gbps connection)!

Solution:

- • Compression (up to 200:1)
 - Using empirical knowledge of human vision & video streams
 - Lossy compression (cf. lossless compression)
- • Reduce frame rate or resolution
- Constrained by practical constraints





How does the problem change with resolution?

If we double the resolution of a 2D image

- We x4 the number of pixels (n^2)
- If we have a 3D image (e.g. from an MRI scanner)
 - If we double the resolution then: x8 (n^3)

The same problem will occur with 2D models and will be even worse as the number of dimensions increases:

- n^2 , n^3



Time complexity

To develop a *basic* understanding
of algorithm *design* and *efficiency* in terms of
worst case time complexity.




Outline / Topics

- ◆ Algorithm Design and Analysis
- ◆ Efficiency
- ◆ Time Complexity
- ◆ Big-O Notation
- ◆ Examples
- ◆ Lessons




What is an Algorithm?

In mathematics, computer science, and related subjects, an algorithm is an effective method for solving a problem expressed as a finite sequence of instructions.




- ○ Abu Abdullah Muhammad bin Musa Al-Khwarizmi
 - a Muslim mathematician
 - invented *algorithms* to solve quadratic equations
 - the word *algorithm* derived from his Latin name
Algorithmi
- An even earlier *algorithm* is the sieve of Eratosthenes



Algorithm: question?

Which one is an algorithm?

- ↳ ♦ A recipe for making tomato soup?
- ↳ ♦ A procedure to sort 1000 numbers into ascending numeric order?
- ↳ ♦ A procedure to recognize a particular face in a crowd?
- ♦ A method to order objects according to beauty?



Algorithm Design and Analysis

→ ◆ Multiple algorithms often exist for the same task:

→ ■ All of them give correct results

→ ■ How do we select the best one?

◆ Many possible (and often conflicting) criteria:

■ Efficiency

■ ?

■ simplicity, clarity


■ elegance, proofs of correctness

◆ We need to ask

■ is my algorithm correct? ✓

■ does my algorithm always terminate? ✓

■ does an algorithm even exist? ✓




Algorithm Efficiency

◆ Resource usage of an algorithm

- typically: time (runtime) and space (computer memory)
- also: network usage, hardware requirements, ...
- consider trade-offs between resources

↳ ◆ How do we measure the run-time of an algorithm?

- benchmarking on *representative* set of inputs: empirical analysis
- analyse the (time) complexity



Algorithm Efficiency – Empirical Analysis

◆ Idea: Implement the algorithm (a program) and time it



Question: How can we time a program?

- ■ Manual: Using stopwatch
- ■ Automatic: Using some timer function

Run the program for various input sizes and measure run time.

data size	time (ms)
→ 250	5
500	8
1000	10
2000	15
4000	30
<u>8000</u>	<u>50</u>
16000	75

100,000



Algorithm Efficiency – Time Complexity

◆ Time complexity:

- the number of operations that an **algorithm** requires to execute, in terms of the **size** of the input or problem


Note:

- algorithm, not implementation
 - so: pseudocode; no fixed programming language, computer architecture
- "in terms of" – complexity defined as a function **$T(n)$**

size
of input

◆ Questions:

- ■ what do we mean by *operations*?
- ■ what do we mean by *size*?
- ■ we usually focus on **worst-case**, not **average-case**, analysis



Example # 1

- ◆ Look up a value v in an array x of integers

→


1	4	17	3	90	79	4	6	81
---	---	----	---	----	----	---	---	----

- ◆ Algorithm: Linear Search

- inputs: array x of size n , integer v
- return: index of first occurrence of v in x , or -1 if none

- ◆ $T(n) = n$

```
for i=0...n-1:  
    if x[i] == v:  
        return i  
return -1
```



Example # 2

◆ Matrix-Vector Multiplication: $x = A b$

■ $n \times n$ matrix A , vector of b size n

◆ Algorithm

■ inputs: matrix A , vector b


■ result stored in vector x (initially all 0)

```
for i=0...n-1:
    for j=0...n-1:
        x[i] = x[i] + A[i][j] * b[j]
```

Handwritten annotations: Blue arrows point from the first 'n' to the 'n-1' in the first loop, and from the second 'n' to the 'n-1' in the second loop. A circled '1' is next to the first loop, and a circled '2' is next to the second loop. The assignment statement is underlined.

◆ $T(n) = 2n^2$

$n \times n \times 2$




Quadratic Growth

Consider the two functions

→ $f(n) = n^2$


→ $g(n) = n^2 - 3n + 2$

Around $n = 3$, they look very different




Quadratic Growth

Around $n = 10$, they look similar




Quadratic Growth

Around $n = 100$, they are almost the same



Quadratic Growth

Around $n = 1000$, the difference is indistinguishable!



Asymptotic

Quadratic Growth

The absolute difference is large, such that,

$$\begin{aligned}\underline{f(1000)} &= \underline{1000000} \\ \underline{g(1000)} &= \underline{997002}\end{aligned}$$

but the relative difference is very small,

$$\left| \frac{f(1000) - g(1000)}{f(1000)} \right| = 0.002998 < 0.3\%$$

and this difference goes to zero as $n \rightarrow \infty$



Big-O Notation

◆ Usually we don't need exact complexity $T(n)$


-
- it suffices to know the complexity class
 - we ignore constant factors/overheads, lower orders
 - focus on performance for large n ("asymptotic")

◆ Big-O notation

- for example: $O(n^2)$ "O of n squared" "on the order of n^2 "

◆ Examples

- $T(n) = n \Rightarrow$ complexity class = $O(n)$
- $T(n) = n+2 \Rightarrow$ complexity class = $O(n)$
- $T(n) = 2n^2 \Rightarrow$ complexity class = $O(n^2)$



Big-O Notation

◆ More Examples

- $T(n) = 10n^3 + 1 \Rightarrow$ complexity class = $O(n^3)$
- $T(n) = 5(n+2) \Rightarrow$ complexity class = $O(n)$
- $T(n) = 1000 \Rightarrow$ complexity class = $O(1)$
- $T(n) = n^2 + n + 1 \Rightarrow$ complexity class = $O(n^2)$

constant

◆ Determining the complexity class


- intuitively, it suffices to count the number of loops and the number of times they are executed



Big-O Notation: Common Complexity Classes $n = 64$

◆ Some common complexity classes:

- $O(1)$ = “Constant”
- $O(\log_2 n)$ = “Logarithmic”
- $O(n)$ = “Linear”
- $O(n \log_2 n)$ = “Log Linear”
- $O(n^2)$ = “Quadratic”
- $O(n^3)$ = “Cubic”
- $O(2^n)$ = “Exponential”



◆ Polynomial: $O(n)$, $O(n^2)$, $O(n^3)$, ... – “tractable”

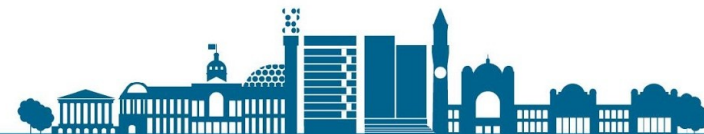
◆ Exponential: $O(2^n)$, $O(C^n)$, – “intractable”



Some concrete numbers ...

- ◆ How many operations needed for an algorithm?
 - with complexity $T(n) = f(n)$ and input size n

$f(n)$	$n = 4$	$n = 16$	$n = 256$	$n = 1024$	$n = 1048576$
1	1	1	1	1.00×10^0	1.00×10^0
$\log_2 \log_2 n$	1	2	3	3.32×10^0	4.32×10^0
$\log_2 n$	2	4	8	1.00×10^1	2.00×10^1
n	4	16	256	1.02×10^3	1.05×10^6
$n \log_2 n$	8	64	2.05×10^3	1.02×10^4	2.10×10^7
n^2	16	256	6.55×10^4	1.05×10^6	1.10×10^{12}
n^3	64	4.10×10^3	1.68×10^7	1.07×10^9	1.15×10^{18}
2^n	16	65536	1.16×10^{77}	1.80×10^{308}	6.74×10^{315652}

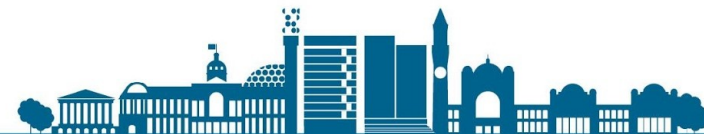


Some concrete numbers ...


◆ How much **time** needed for an algorithm?

■ assuming 1 million operations per second

$f(n)$	$n = 4$	$n = 16$	$n = 256$	$n = 1024$	$n = 1048576$
1	1 usec	1 usec	1 usec	1 usec	1 usec
$\log_2 \log_2 n$	1 usec	2 usec	3 usec	3.32 usec	4.32 usec
$\log_2 n$	2 usec	4 usec	8 usec	10 usec	20 usec
n	4 usec	16 usec	256 usec	1.02 msec	1.05 sec
$n \log_2 n$	8 usec	64 usec	2.05 msec	10.2 msec	21 sec
n^2	16 usec	256 usec	65.5 msec	1.05 sec	12.72 days
n^3	64 usec	4.1 msec	16.8 sec	17.83 min	36559 yrs
2^n	16 usec	65.5 msec	3.7×10^{63} yrs	5.7×10^{294} yrs	2.14×10^{315639} yrs



Graphically



Example # 1 (again)

$$v = \underline{90}$$

- ◆ Look up a value v in an array x

1	4	17	3	90	79	4	6	81
---	---	----	---	----	----	---	---	----



Example # 1 (again)

- ◆ Look up a value v in a sorted array x

1	3	4	4	6	17	79	81	90
---	---	---	---	---	----	----	----	----



$$v = 81$$

Example # 1 (again)

- ◆ Look up a value v in a sorted array x


1	3	4	4	6	17	79	81	90
---	---	---	---	---	----	----	----	----

- ◆ Algorithm: Binary Search

- inputs: sorted array x of size n , integer v
- return: index of first occurrence of v in x , or -1 if none

- ◆ Complexity = $O(\log_2 n)$

```
left = 0; right = n-1
while left < right:
    mid = (left+right)/2
    if x[mid] < v:
        left = mid+1
    else:
        right = mid
if x[left] == v:
    return left
else:
    return -1
```



Computing Complexity Classes

- ◆ Determine total algorithm complexity
 - from the complexities of its components

- 1) Sequential algorithm phases
- 2) Function / Method calls



$$\max\left(\frac{O(n)}{1} + \frac{O(n^2)}{1}\right)$$


Sequential Algorithm Phases

- ◆ Example: matrix-vector multiplication: $x = A b$
 - matrix-vector multiplication $x = A b$
 - initialise vector x (all 0), then add values to x

```
for i=0...n-1:
    x[i] = 0
    for j=0...n-1:
        x[i] = x[i] + A[i][j] * b[j]
```

Handwritten annotations: $O(n)$ for the first loop, n for the second loop, and $O(n^2)$ for the inner loop.

- ◆ Complexity: $O(n) + O(n^2) = O(n^2)$
 - In general: “maximum” of complexities




$$O(n \log_2 n)$$

Function/Method calls

- ◆ Example: n array look-ups


$$O(n)$$



```
for i=0...n-1:  
    → binary_search(x, vi)
```

$$O(\log_2 n)$$

- ◆ Complexity: $O(n) \times O(\log n) = O(n \log n)$
- ◆ In general: “multiplication” of complexities



Computing Complexity Classes

- ◆ Determine total algorithm complexity
 - from the complexities of its components


1) Sequential algorithm phases: “maximum”



- e.g. $O(n) + O(n^2) = O(n^2)$
- e.g. $O(n) + O(\log n) = O(n)$

2) Function / Method calls: “multiplication”

- e.g. $O(n) \times O(\log n) = O(n \log n)$
- e.g. $O(n^2) \times O(1) = O(n^2)$

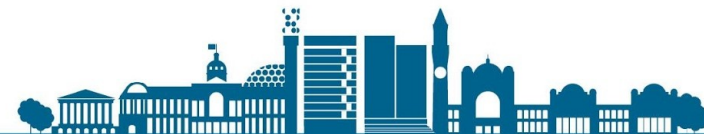


Question # 1

Given the following code fragment, what is its Big-O running time?

```
1  test = 0
2  for i in range(n):
3      for j in range(n):
4          test = test + i * j
```

- a) $O(n)$
- b) $O(n^2)$
- c) $O(\log n)$
- d) $O(n^3)$



Question # 2

Given the following code fragment, what is its Big-O running time?

```
1  test = 0
2  for i in range(n):
3      test = test + 1
4
5  for j in range(n):
6      test = test - 1
```

- a) $O(n)$
b) $O(n^2)$
c) $O(\log n)$
d) $O(n^3)$



Question # 3

Given the following code fragment, what is its Big-O running time?

```
1 i = n
2 while i > 0:
3     k = 2 + 2
4     i = i / 2
5 }
```

1500

1000


500

250

125

...

- a) $O(n)$
- b) $O(n^2)$
- c) $O(\log n)$
- d) $O(n^3)$




$$9! = 9 \times 8 \times 7 \times \dots \times 1$$

Some Harder Problems

◆ Traveling Salesman Problem (TSP)


- given n cities and the distances between them, what is the **shortest possible route** that **visits each city exactly once** and then **returns to the first city**?



“A naive approach to solving TSP would be the brute-force solution, which means finding all possible routes given a certain number of nodes. This is a very expensive way to solve it, with a time complexity of $O(n!)$.”

<https://tonicanada.medium.com/introduction-to-the-travelling-salesman-problem-5ace44932cb5>

<https://www.youtube.com/watch?v=SC5CX8drAtU>



Some Harder Problems




◆ Boolean Satisfiability Problem (SAT)

- Given a Boolean formula $F(x_1, x_2, x_3, \dots, x_n)$
- Can F evaluate to 1 (true)?
- If yes, return values of x_i 's (satisfying the assignment) that make F true.

◆ In both cases (TSP & SAT):

- lots of practical applications (Operations Research, Optimization, Logistics, Model Checking, Software Verification, ... etc.)
- also important problems in theoretical computer science




Algorithms vs. Problems

◆ Algorithm complexity

- worst-case run-time of algorithm – "efficiency"
- actually an upper bound: algorithm A is in $O(f(n))$ if the worst-case run-time is at most $f(n)$
- usually use tightest (most informative) complexity

◆ Problem complexity

- complexity class = set of problems
- problem X is in complexity class $O(f(n))$ if there exists an algorithm to solve it in $O(f(n))$ – "difficulty"
- again: this is an upper bound (and uses the tightest possible)
- sometimes consider lower bounds:
- e.g. sorting: $O(n \log n)$



P and NP

◆ Complexity classes

- $O(1) \subseteq O(\log n) \subseteq \underline{O(n)} \subseteq \underline{O(n \log n)} \subseteq \underline{O(n^2)} \subseteq \underline{O(2^n)}$
- polynomial time (PTIME or P) – assumed to be "tractable"

◆ Another famous class: NP (Non-deterministic polynomial time)

- if we can "guess" a solution, it can be checked efficiently (efficiently = in polynomial time). e.g. Sudoku, Chess (more [at this link](#))
- In contrast, polynomial time (or P) problems are those where finding the answer is easy.


◆ NP-hard problems

- e.g. travelling salesman problem, SAT, ...
- only exponential time algorithms are known
- but some efficient heuristics exist in practice


◆ P = NP?

◆ Nobody knows ...

◆ [US\\$1,000,000 prize if you solve it!](#)



				3	7	6		
			6				9	
		8						4
	9							1
6								9
3							4	
7						8		
	1				9			
		2	5	4				



Summary

◆ Algorithm design and analysis:

→ ■ Efficiency

◆ Time Complexity

→ ■ (worst-case) number of operations an algorithm needs to execute, in terms of the size of the input or problem: $T(n)$

◆ Big-O notation


■ complexity classes: $O(n)$, $O(n \log n)$, $O(2^n)$, ...

■ focus on large values of n (i.e. asymptotic behaviour)

■ ignore constants, lower factors

■ count loops/iterations, decompose algorithm

→ ■ complexities for algorithms/problems



Lessons

- ◆ Data Representation/Algorithms are tightly linked
- ◆ Some algorithms are just better than others:
 - ◆ But sometimes the ‘best’ depends on the problem
- ◆ Some problems are inherently hard (in the complexity sense)
 - ◆ This doesn’t mean they cannot be solved
 - ◆ We may need to reformulate the problem
 - ◆ Constrain it
 - ◆ Relax conditions
 - ◆ E.g. instead of “find the best” we “find a very good”



Books / References

- ◆ Goldschlager and Lister: Computer Science: a Modern Introduction (Prentice Hall, 2nd edition 1988)
 - Chapter 3
- ◆ Aho & Ullman: Foundations of Computer Science (Freeman, 1992)
 - Chapter 3
 - <http://infolab.stanford.edu/~ullman/focs.html>

