



UNIVERSITY OF
BIRMINGHAM

Computer Systems Process Scheduling



Lecture Outline

- ◆ Basic Concepts of CPU Scheduling
- ◆ Scheduling Criteria
- ◆ Scheduling Algorithms
 - FCFS – First Come First Serve
 - SJF – Shortest Job First
 - SRTF – Shortest Remaining Time First
 - Priority and Round-Robin Scheduling



Basic Concepts

- ◆ **Single processor** system = **single process** can run at a time
- ◆ Others must wait until the CPU is free and can be rescheduled
- ◆ Objective of multi-programming is to have **some process running at all times**
 - **Maximize CPU utilization**
- ◆ This is a fundamental OS function

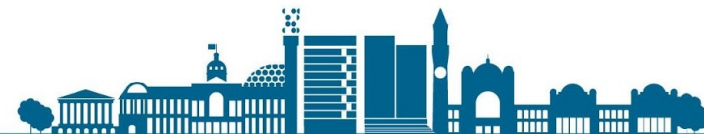
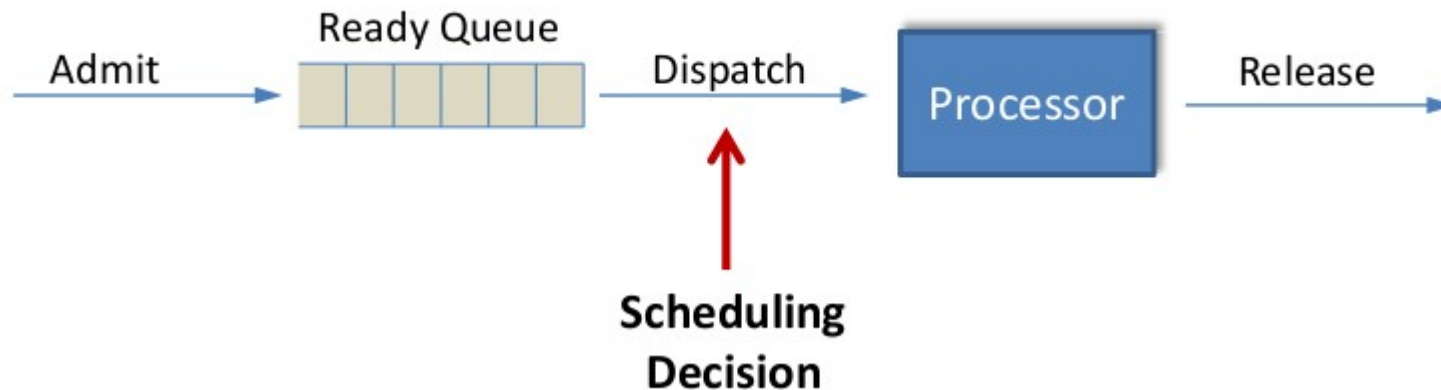
Every OS has to have a scheduler.



Scheduling

◆ CPU Scheduling


- Decides which process to execute next – is the activity of **selecting the next process** in the Ready queue for execution on a processor



When does Scheduling happen?

- ◆ A **scheduling decision** takes place when an event occurs that interrupts the execution of a process
- ◆ Possible events
 - Clock Interrupts (i.e. process **running** → **ready** state)
 - I/O Interrupts (i.e. process **waiting** → **ready** state)
 - Operating System Calls (read, write, process **ready** → **waiting**)
 - Signals (e.g. Semaphores)

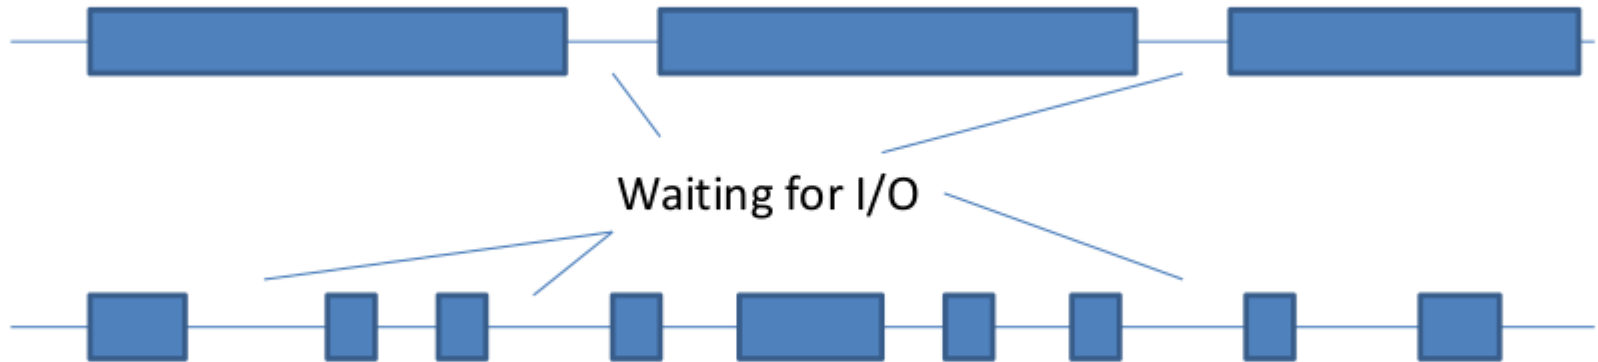
A lock




Characteristics of Process Execution

- ◆ Batch Processing: **CPU Bound**
- ◆ Interactive Systems: **I/O Bound**

CPU – intensive Execution: long CPU bursts

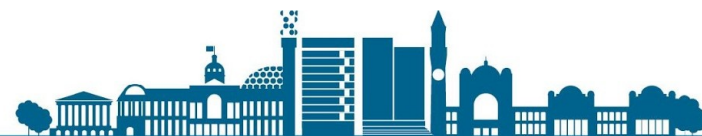
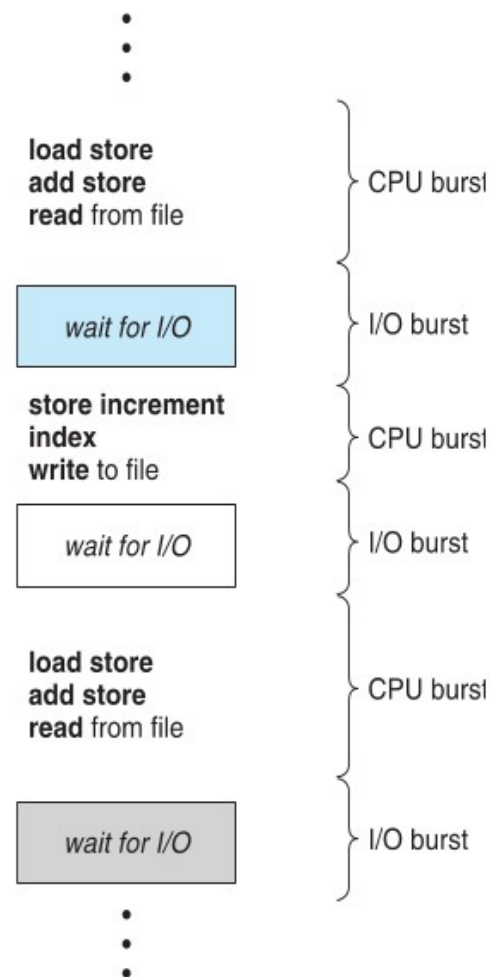


I/O – intensive Execution: short CPU bursts




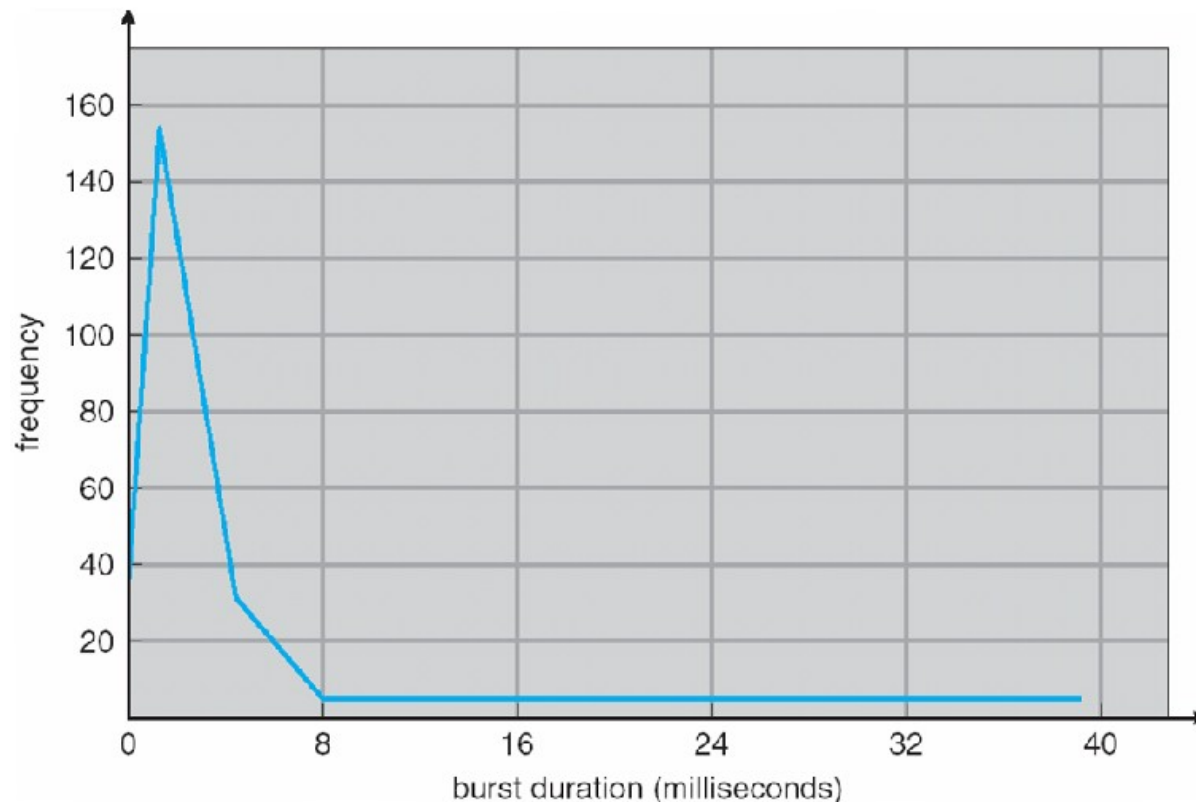
CPU-I/O Burst Cycle

- ◆ Process execution consists of a **cycle** of CPU execution and I/O wait
- ◆ Processes **alternate** between these two states
- ◆ Process execution begins with a **CPU burst**, this is followed by an **I/O burst**. This pattern repeats!
- ◆ Scheduler **can schedule another process** during I/O burst.
- ◆ Eventually **final CPU burst** ends with a system request to terminate execution



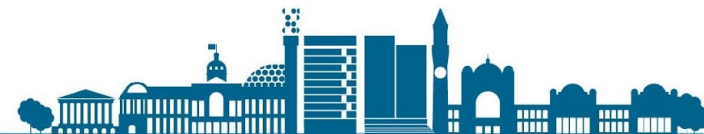
Histogram of CPU Bursts

- ◆ This histogram shows that most of the processes have many **very short CPU bursts** e.g. Interactive systems



CPU Scheduler

- ◆ Recall: whenever the CPU becomes idle, the OS must select one of the processes in the ready queue to be executed next
 - This is carried out by the **short-term scheduler** from the processes in the ready queue
- ◆ The queue is **not necessarily** first-in, first-out (FIFO)
 - It may be a priority queue, tree or an unordered linked list
- ◆ Conceptually, all processes are **waiting for a chance** to run on the CPU

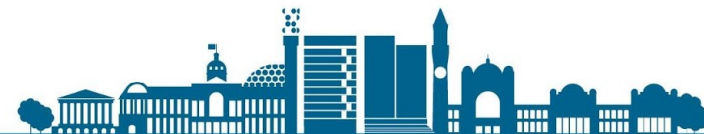


Dispatcher

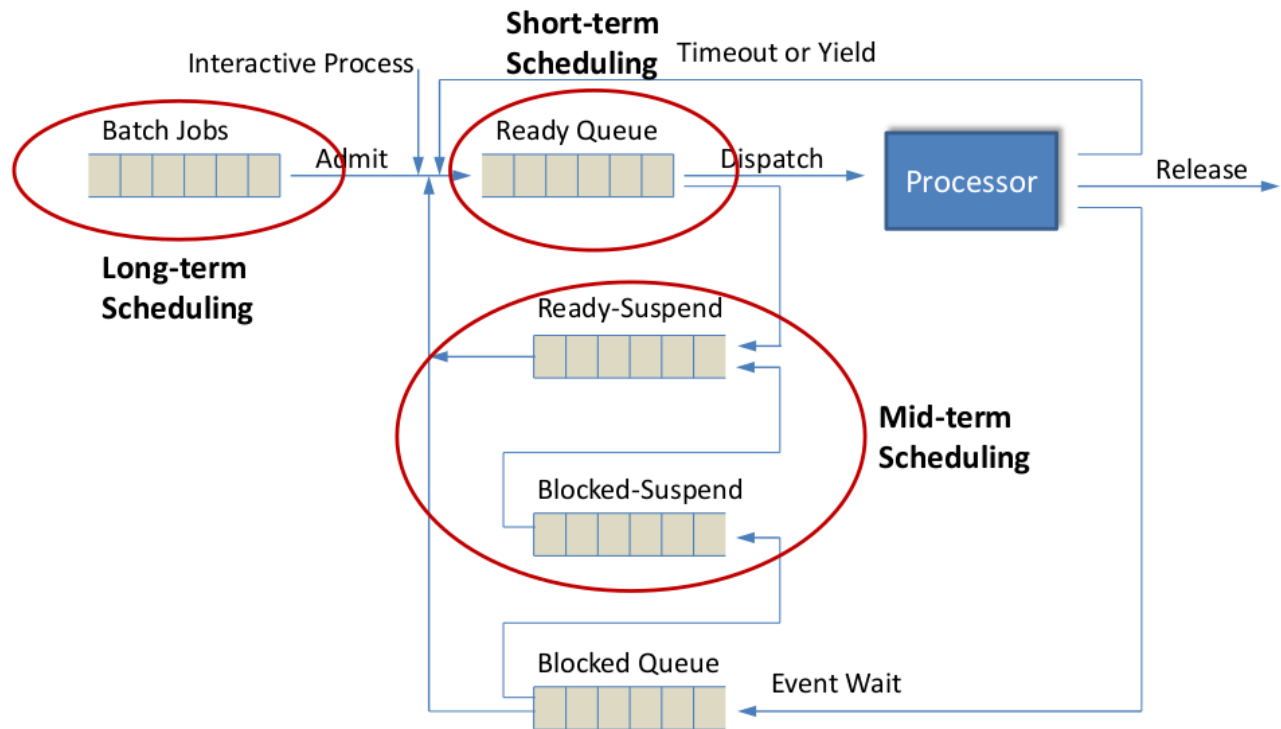
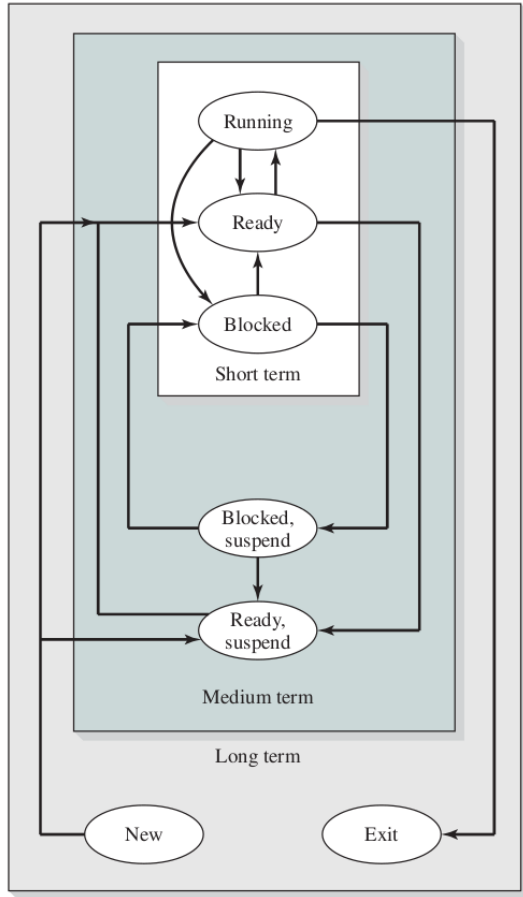
- ◆ The **dispatcher** is the module that **gives control of the CPU** to the process selected by the short-term scheduler. This involves:
 - Switching context
 - Switching to user mode
 - Jumping to the proper location in the user program to restart that program
- ◆ Dispatcher should be as fast as possible
 - It is invoked during **every process switch**
 - Time taken by dispatcher to stop one process and start another one is known as **dispatch latency**.

调度员

If dispatcher is slow, this very reduce the speed
of a computer significantly

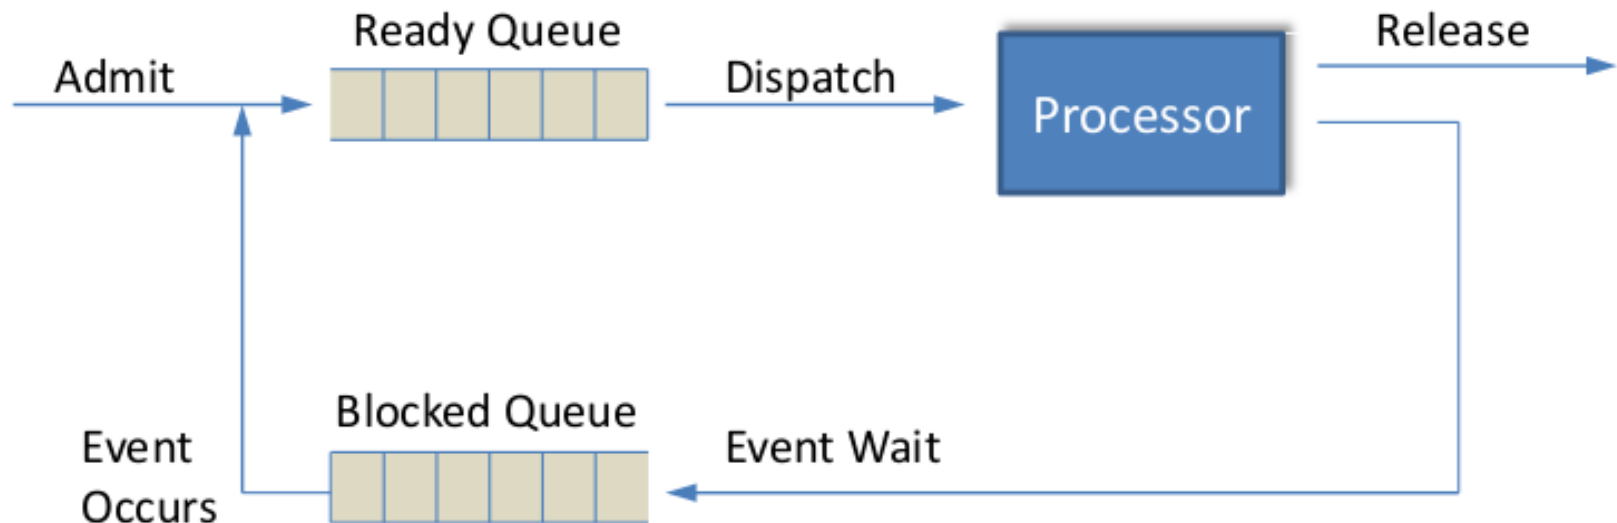


Levels of Scheduling



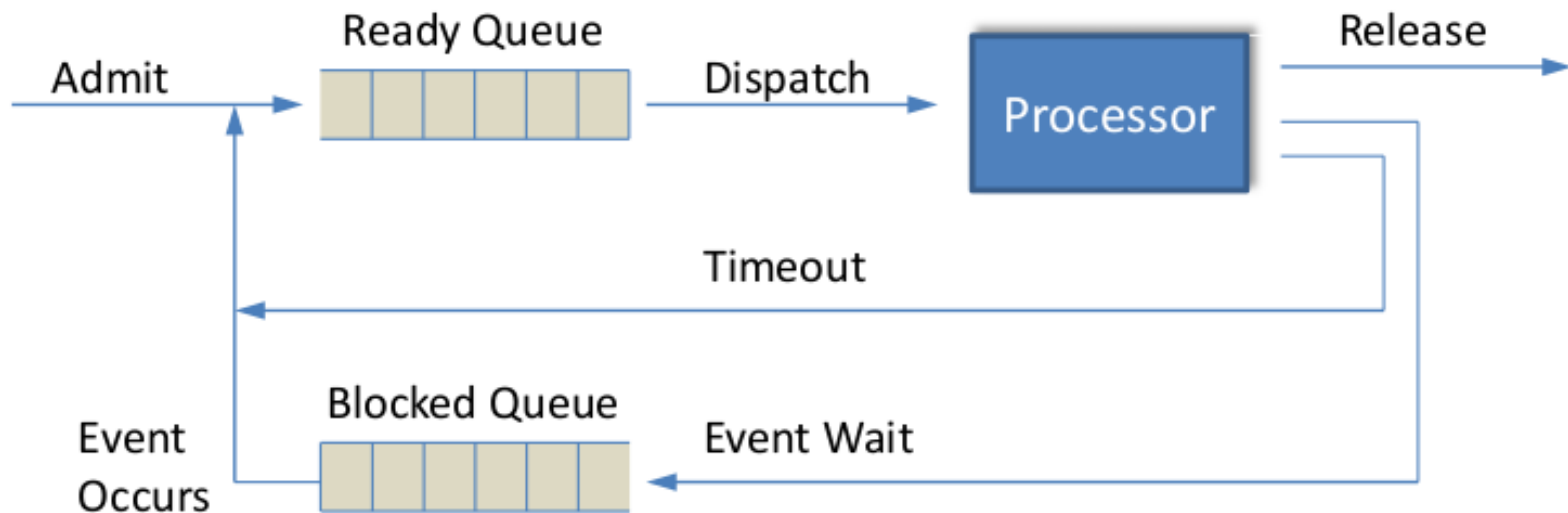
Non-preemptive Scheduling

- ◆ Once a process is scheduled, it continues to execute on the CPU, until:
 - it is **finished** (terminates)
 - it **releases** the CPU **voluntarily** (cooperative scheduling)
 - It **blocks** due to an event such as I/O interrupts, waits for another process etc.




Preemptive Scheduling

- ◆ The operating system interrupts (intervenes) processes
 - A scheduled process executes, until its time slice is used up
 - Clock interrupt returns control of CPU to scheduler (time slice expired)
 - ▶ Current process is **suspended** and placed in the Ready queue
 - ▶ New process is selected from Ready queue and executed on CPU
 - When a process with **higher priority** becomes ready



Scheduling Criteria

- ◆ Different CPU-scheduling algorithms have different properties. One may favour one class of processes over another.
- ◆ Many criteria suggested, including
 - **CPU Utilization** - keep that CPU as busy as possible. 0-100%, but realistically, 40-90%
 - **Throughput** - Number of processes that complete their execution per time unit
 - **Turnaround time** - Total amount of time to execute one process to its completion.
 - **Waiting time** - Total amount of time a process has been waiting in the Ready queue
 - **Response time** - The time it takes from when a request was submitted until the first response is produced by the operating system (latency).
 - Others include meeting deadlines, predictability, fairness, balance and enforcing priorities.



Optimization Criteria

◆ Maximize (operating system concerns)

- CPU Utilization
- Throughput

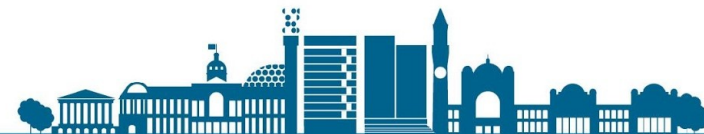
◆ Minimize (user concerns)

- Turnaround time
- Waiting time
- Response time



First Come, First Served Scheduling (FCFS)

- ◆ Simplest CPU-scheduling algorithm
- ◆ The process that **requests the CPU first** is allocated the CPU first
- ◆ Managed with a FIFO queue, process enters the ready queue
- ◆ PCB linked onto tail of queue
- ◆ When CPU free, it is allocated to the process at the head of the queue
- ◆ Running process then removed from queue




FCFS Example

- ◆ Suppose the processes arrive in the order: P₁ , P₂ , P₃. Gantt Chart as follows:

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3

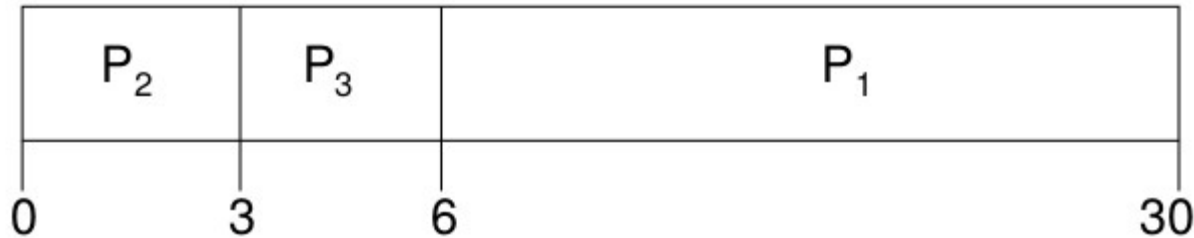


- ◆ Waiting time for P₁ = 0; P₂ = 24; P₃ = 27
- ◆ Average waiting time: $(0 + 24 + 27)/3 = 17\text{ms}$



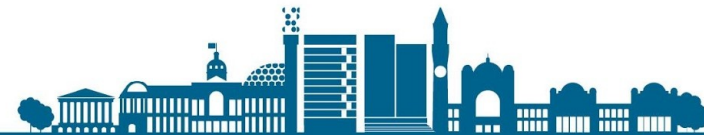
FCFS Example - 2

- ◆ Suppose now that the processes arrive in the order: P_2 , P_3 , P_1 . Gantt Chart for schedule:




Process	Burst Time
P ₁	24
P ₂	3
P ₃	3

- ◆ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- ◆ Average waiting time: $(6 + 0 + 3)/3 = 3\text{ms}$
- ◆ Better, but shows that waiting time under an FCFS policy may vary substantially



Shortest-Job First (SJF) Scheduling


- ◆ Associate with each process the **length** of the process's **next CPU burst**
- ◆ When the CPU is available it is assigned to the process that has the **smallest next CPU burst**
 - If next CPU bursts of two processes are the same, FCFS is used to break the tie
- ◆ Better name **shortest-next-CPU-burst** algorithm - scheduling depends on the length of the next CPU burst of a process, rather than total length



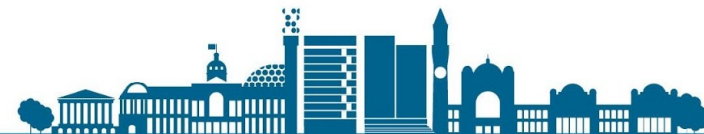
SJF Example

<u>Process</u>	<u>Burst Time</u>
P ₁	6
P ₂	8
P ₃	7
P ₄	3

SJF scheduling chart



- ◆ Average waiting time: $(3 + 16 + 9 + 0)/4 = 7\text{ms}$
- ◆ In comparison, FCFS would be 10.25ms



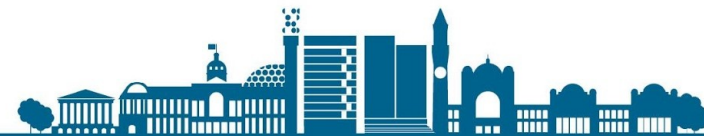
Optimality of SJF

- ◆ Provably optimal - it gives the **minimum average waiting time** for a given set of processes
- ◆ Moving a short process before a long one **decreases the waiting time** of the short process **more than it increases the waiting time** of the long process - hence average time decreases



Difficulty of SJF

- ◆ Determining the length of the next CPU request is difficult!
- ◆ Could ask the user (if they know)?
- ◆ For this reason, SJF is **difficult to implement** at the level of short-term CPU scheduling
- ◆ There is **no way to know** the length of the next CPU burst
- ◆ Approximation is a possible solution



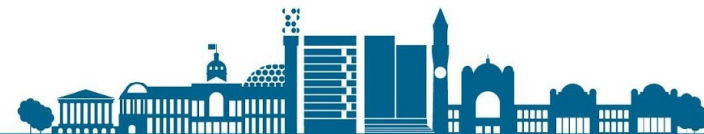
Determining the Length of the next CPU burst

- ◆ **Exponential averaging** of the measured lengths of previous CPU bursts:

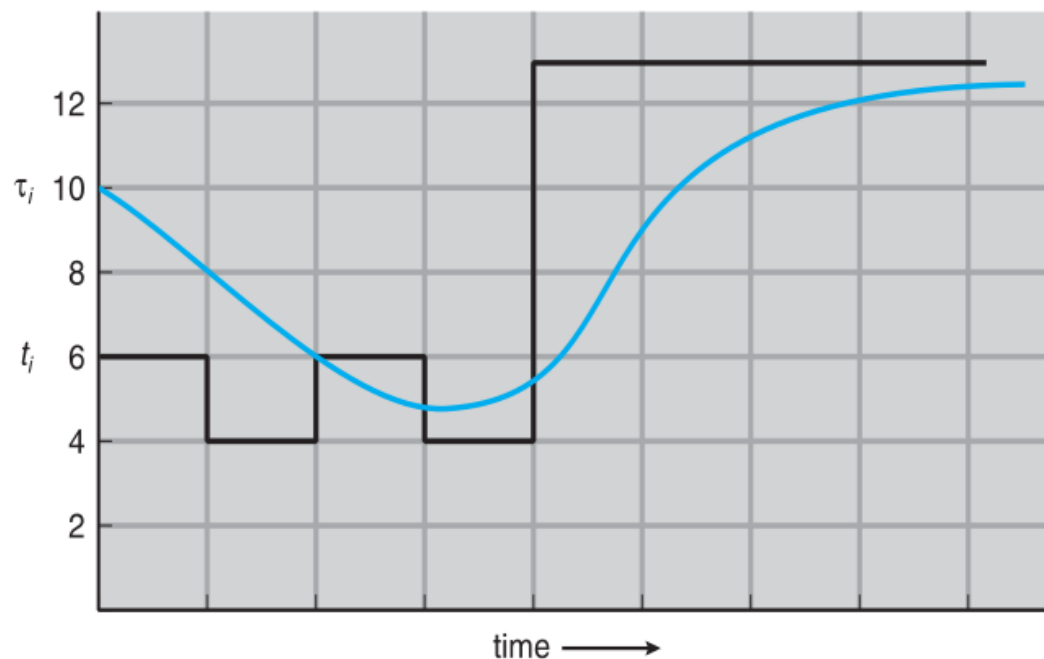
1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$$

Commonly, α set to 0.5



Prediction of the Length of the next CPU Burst



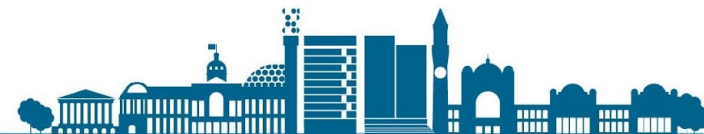
CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n$$

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0.$$

Shortest Remaining Time First Scheduling

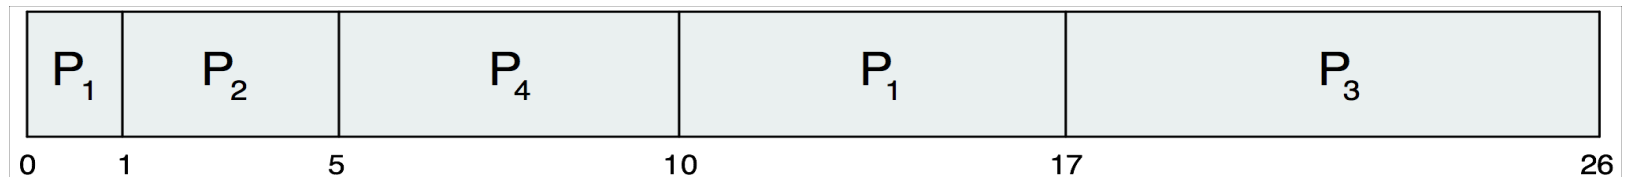
- ◆ SJF is either preemptive or nonpreemptive
- ◆ Choice arises when a **new process arrives** at the ready queue while a **previous process still executing**
- ◆ **Next CPU burst** of the newly arrived process **may be shorter** than what is left of the currently executing process
- ◆ Preemptive SJF will preempt the currently executing process
- ◆ Nonpreemptive will allow the currently running process to finish its CPU burst
- ◆ Preemptive SJF scheduling called **Shortest-Remaining-Time-First (SRTF)** scheduling




SRTF Scheduling – Example

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P ₁	0	8
P ₂	1	4
P ₃	2	9
P ₄	3	5

◆ Preemptive SJF Gantt Chart:



◆ Average waiting time = $((10-1)+(1-1)+(17-2)+(5-3))/4$
 $= 26/4 = 6.5\text{ms}$



Priority Scheduling


- ◆ SJF special case of general **priority-scheduling** algorithm
- ◆ Priority associated with each process
- ◆ CPU allocated to process with **highest priority**
 - Equal priority scheduled using FCFS



Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

- ◆ Low vs. high priority
- ◆ Priorities indicated by fixed range of numbers
- ◆ **Low numbers = high priority**



Priority Scheduling


- ◆ Priorities can be defined internally or externally
- ◆ **Internal criteria** = measurable quantity such as time limits, memory requirements, number of open files
- ◆ **External criteria** = Outside OS such as process importance, type and amount of funds paid for computation, department sponsoring the work, *political factors*



Priority Scheduling – Preemption


Preemptive vs. Non-preemptive

- ◆ Preemptive will **preempt** the CPU if the priority of the newly arrived process in the ready queue is **higher than** the priority of the currently running process
- ◆ Non-preemptive priority scheduling algorithm will put the **new process at the head** of the ready queue



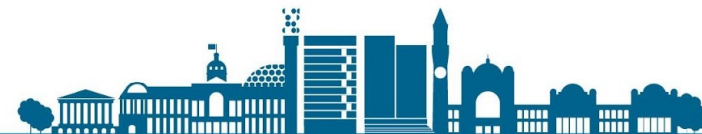
Priority Scheduling – Problems

- ◆ Indefinite **blocking** or **starvation**
- ◆ A process ready to run but waiting for CPU is blocked
- ◆ **Low priority** processes can be **blocked indefinitely**
- ◆ A heavily loaded computer system with a steady stream of high priority processes can prevent a low-priority process from ever getting the CPU
- ◆ Two things will probably happen:
 - Process will **eventually run** (at 2 A.M Sunday)
 - Computer will **eventually crash** and lose these unfinished, low priority processes
- ◆ **Solution:** importance increases with age (**aging**)
 - Increase importance by a point every 15 minutes



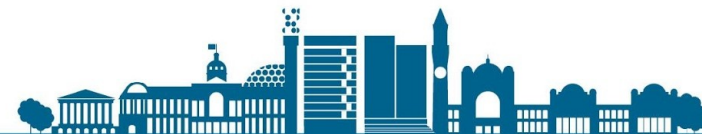
Round-Robin (RR) Scheduling

- ◆ Specially designed for **time-sharing systems**
- ◆ Similar to FCFS, but with preemption to enable the switching of processes
- ◆ Small unit of time used: **time quantum** or **time slice**
 - 10 to 100 milliseconds in length
- ◆ Ready queue is circular
- ◆ CPU goes around the ready queue, allocating CPU time to each process for a time interval or up to 1 time quantum



Round-Robin (RR) Scheduling

- ◆ Implementation treats ready queue as a FIFO queue of processes
- ◆ CPU scheduler **picks first process** from ready queue, **sets a timer** to interrupt after 1 time quantum, and dispatches the process
- ◆ Two things could happen:
 - Process may have CPU burst of **less than 1 quantum**, so process will release CPU **voluntarily**
 - If process CPU burst greater than 1 time quantum, timer goes off and causes an **interrupt to the OS**, resulting in a **context switch** and process goes to the tail of ready queue

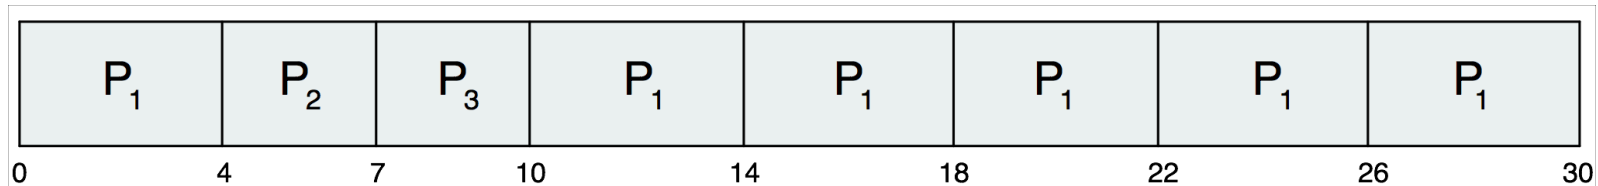


Round-Robin (RR) Scheduling – Example

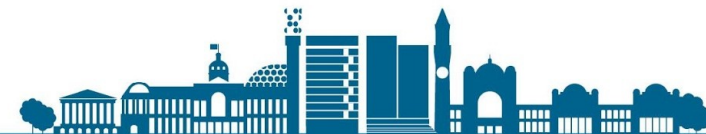
- ◆ Time quantum = 4 milliseconds

<u>Process</u>	<u>Burst Time</u>
P ₁	24
P ₂	3
P ₃	3

Resulting RR schedule as follows:




$$\begin{aligned}\text{Average waiting time} &= ((10 - 4) + 4 + 7)/3 \\ &= 17/3 = 5.66\text{ms}\end{aligned}$$



Round-Robin (RR) Scheduling

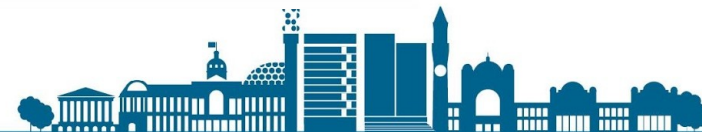
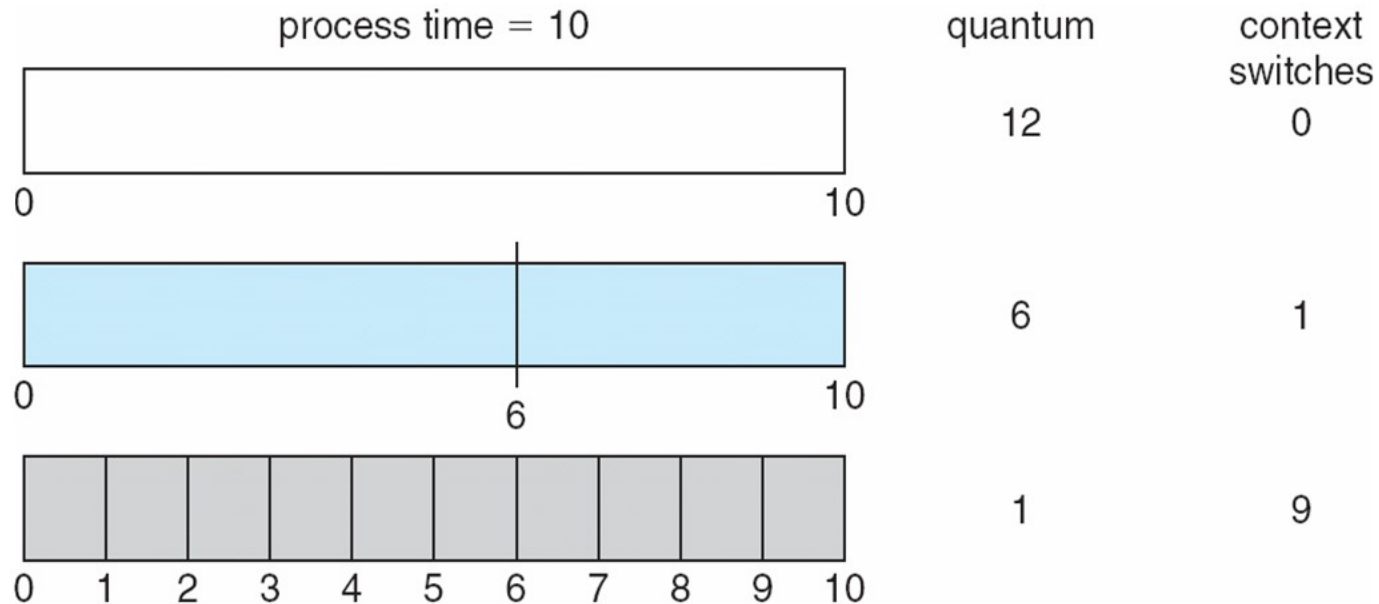
- ◆ If there are n processes in the ready queue and the time quantum is q , each process gets $1/n$ of the CPU time in chunks of at most q time units
- ◆ Each process must wait no longer than $(n - 1) \times q$ time units until its next time quantum
- ◆ Example:
 - 5 processes
 - $q = 20$ milliseconds
 - Each process gets up to 20 milliseconds every 100 milliseconds



RR Scheduling Performance

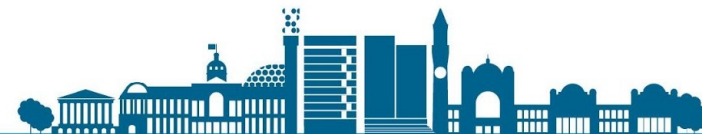
◆ Depends heavily on size of time quantum

- If extremely large, RR is same as FCFS policy
- If extremely small, RR can result in a large number of context switches



RR – Time Quantum vs. Context Switch Time

- ◆ Ideally, time quantum should be large with respect to the context-switch time
- ◆ If context-switch is approximately 10% of time quantum, then ~10% of CPU time will be spent context-switching
- ◆ In reality, most modern systems have time quanta ranging from **10 - 100 milliseconds**
 - In comparison, context-switch is typically less than **10 microseconds**



Summary

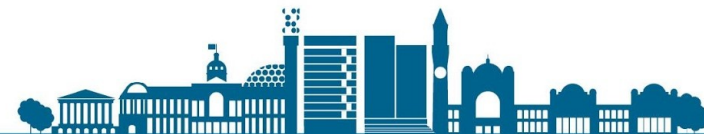
◆ Basics of CPU Scheduling

- CPU-I/O Burst Cycle
- Preemption
- Dispatcher

◆ Scheduling Criteria

◆ Scheduling Algorithms

- FCFS, SJF, SRTF, Priority, Round-Robin



References / Links

- ◆ Chapter # 6: **CPU Scheduling**, Operating System Concepts (9th edition) by Silberschatz, Galvin & Gagne

