




UNIVERSITY OF
BIRMINGHAM

Computer Systems

Week 10b – Network Layer Routing Algorithms

Based on material and slides from
Computer Networking: A Top Down

Approach, 7th Edition – Chapter 5
Jim Kurose, Keith Ross
Pearson/Addison Wesley



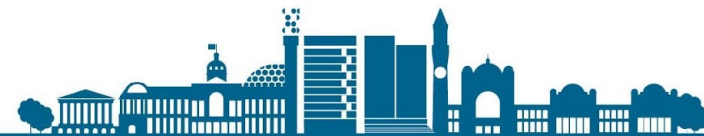
Lecture Objective

The objective of this lecture is to understand the **conceptual** aspects of **network layer** routing algorithms

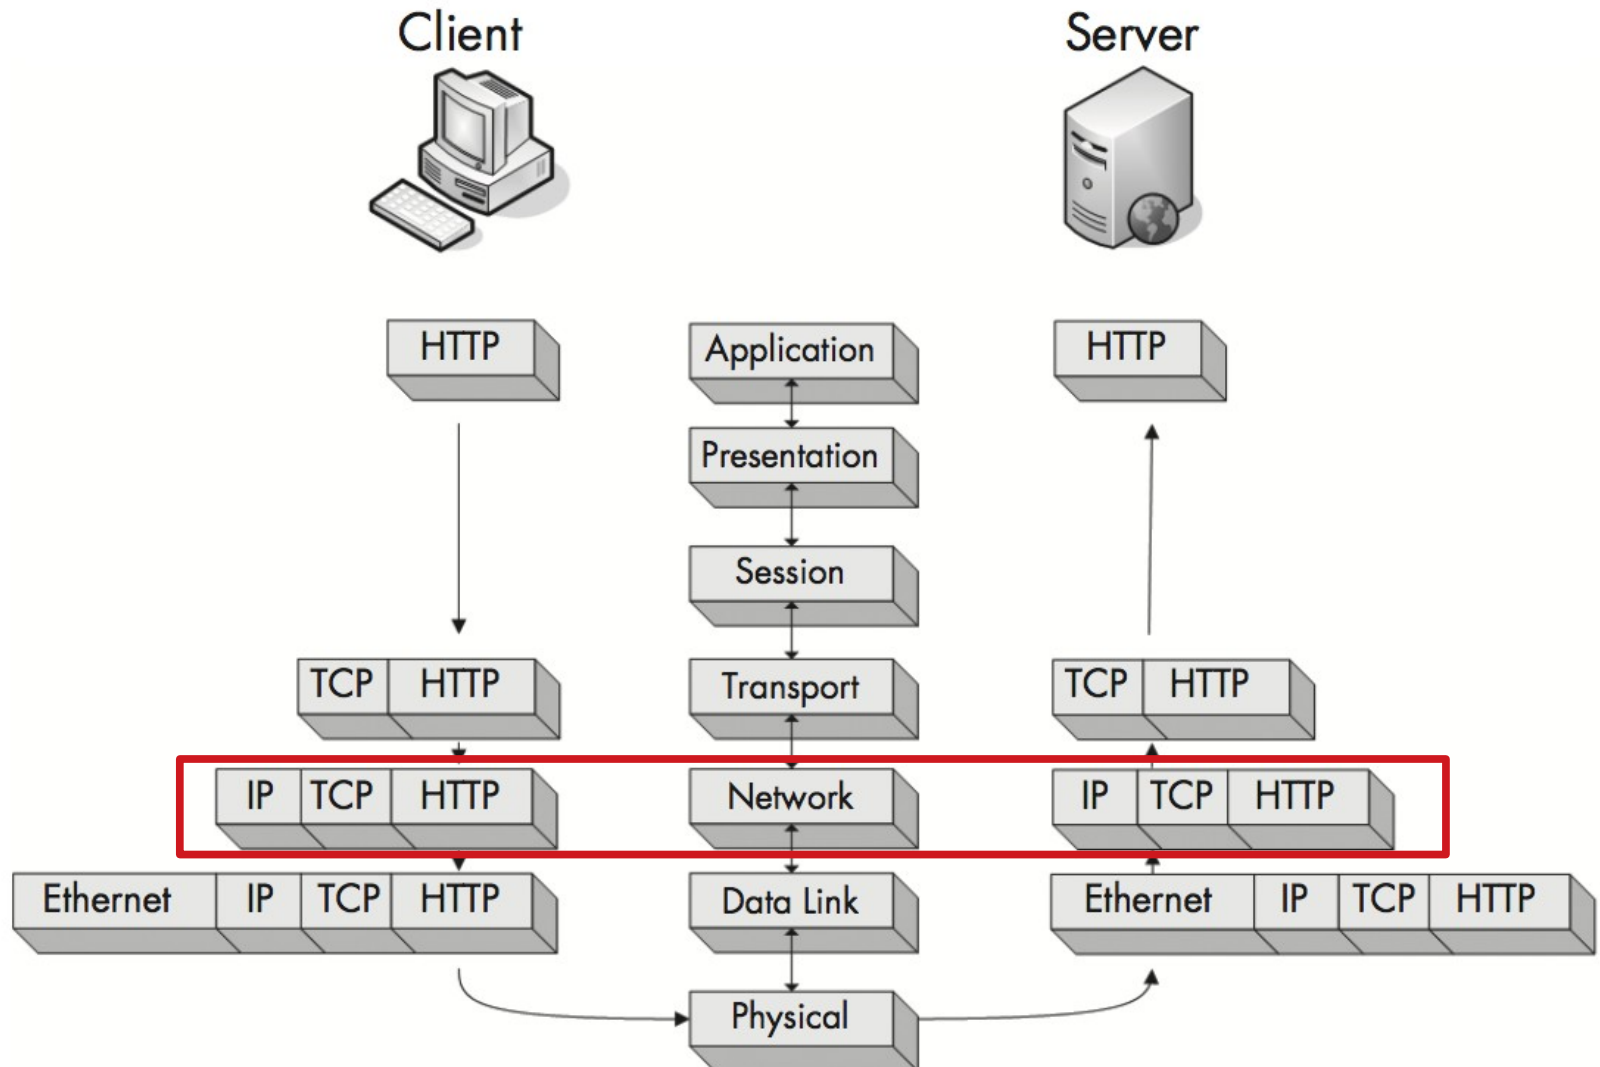


Lecture Outline


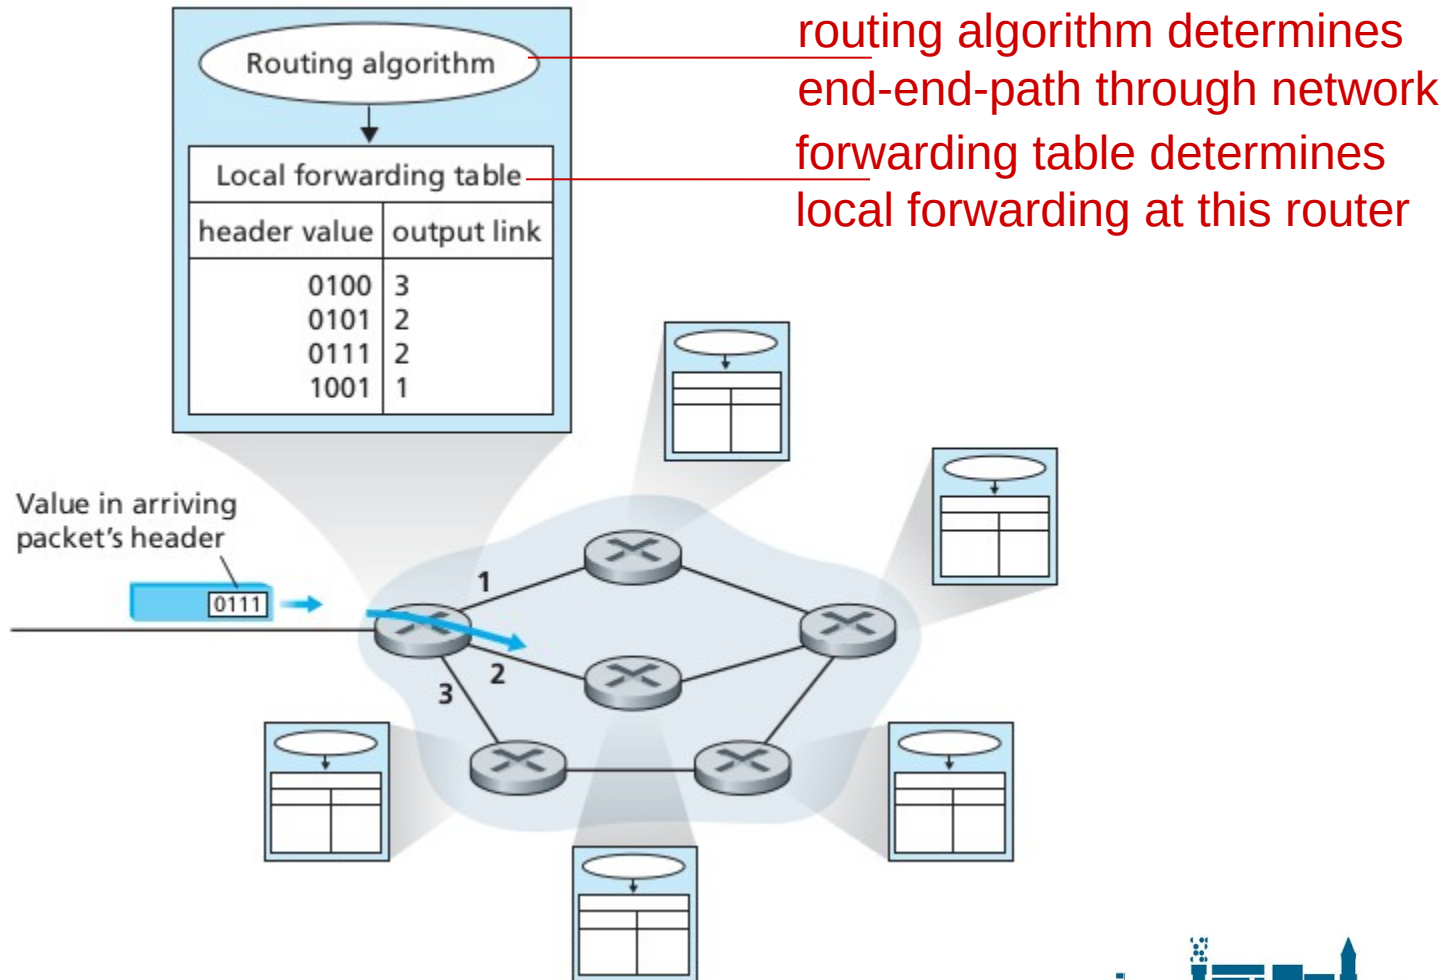
- ◆ Routing vs. Forwarding – Recap
- ◆ Graph Abstraction
- ◆ Routing Algorithms – Classification
- ◆ Link State Routing – Dijkstra's Algorithm
- ◆ Distance Vector Routing – Bellman-Ford Algorithm
- ◆ Hierarchical Routing
- ◆ Summary




Recap - Network Layers



Interplay between Routing & Forwarding



Graph Abstraction




graph: $G = (N, E)$


N = set of routers = $\{ u, v, w, x, y, z \}$

E = set of links = $\{ (u,v), (u,x), (u,w), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Aside: graph abstraction is useful in other network contexts, e.g., P2P, where N is set of peers and E is set of TCP connections



Graph Abstraction



$c(x,x') = \text{cost of link } (x,x')$


e.g., $c(w,z) = 5$

cost could always be 1, or
inversely related to **bandwidth**,
or **inversely** related to
congestion

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Key Question: What is the least-cost path between u and z ?

Routing Algorithm: Algorithm that finds that least cost path



Routing Algorithm Classification

Q: Global or decentralized information?

Global:

- ◆ All routers have complete topology, link cost info
- ◆ “Link State” algorithms

Decentralized:

- ◆ Router knows physically-connected neighbors, link costs to neighbors
- ◆ Iterative process of computation, exchange of info with neighbors
- ◆ “Distance Vector” algorithms

Q: static or dynamic?

Static:

- ◆ Routes change slowly over time

Dynamic:

- ◆ Routes change more quickly
 - periodic update
 - in response to link cost changes




A Link-State Algorithm

Dijkstra's Algorithm

- ◆ Net topology, link costs known to all nodes
 - accomplished via “link state broadcast”
 - all nodes have same info
- ◆ Computes least cost paths from one node (“source”) to all other nodes
 - gives forwarding table for that node
- ◆ Iterative: after k iterations, know least cost path to k destinations

Notation:

- ◆ $c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors
- ◆ $D(v)$: current value of cost of path from source to dest. v
- ◆ $p(v)$: predecessor node along path from source to v
- ◆ N' : set of nodes whose least cost path definitively known



Dijkstra's Algorithm

1 **Initialization:**

2 $N' = \{u\}$

3 for all nodes v

4 if v adjacent to u

5 then $D(v) = c(u,v)$

6 else $D(v) = \infty$

7

8 **Loop**

9 find w not in N' such that $D(w)$ is a minimum

10 add w to N'


11 update $D(v)$ for all v adjacent to w and not in N' :

12 $D(v) = \min(D(v), D(w) + c(w,v))$

13 /* new cost to v is either old cost to v or known

14 shortest path cost to w plus cost from w to v */

15 **until all nodes in N'**




Dijkstra's Algorithm - Example

Step	N'	D(v) p(v)	D(w) p(w)	D(x) p(x)	D(y) p(y)	D(z) p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uwx	6,w			11,w	14,x
3	uwxv				10,v	14,x
4	uwxvy					12,y
5	uwxvyz					


Notes:

- ◆ construct shortest path tree by tracing predecessor nodes
- ◆ ties can exist (can be broken arbitrarily)




Dijkstra's Algorithm - Another Example

Step	N'	D(v),p(v)	D(w),p(w)	D(x),p(x)	D(y),p(y)	D(z),p(z)
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					




Dijkstra's Algorithm - Least Cost Path and Forwarding Table for Node u



Resulting shortest-path tree from u:

Destination	Link
v	(u, v)
w	(u, x)
x	(u, x)
y	(u, x)
z	(u, x)

Resulting forwarding table in u:



Dijkstra's Algorithm - Complexity


Algorithm Complexity: n nodes

- ◆ Each iteration: need to check all nodes, w , not in N'
- ◆ $n(n+1)/2$ comparisons: $O(n^2)$
- ◆ More efficient implementations possible: $O(n \log n)$




Dijkstra's Algorithm - Oscillations


Support link cost equals amount of carried traffic.




a. Initial routing



b. x, y detect better path to w, clockwise




c. x, y, z detect better path to w, counterclockwise



d. x, y, z, detect better path to w, clockwise

Given these costs, find new routing....
resulting in new costs



Distance-Vector (DV) Routing Algorithm

Bellman-Ford equation:

let

$d_x(y) :=$ cost of least-cost path from x to y

then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

cost from neighbor v to destination y


cost to neighbor v

\min taken over all neighbors v of x



Bellman-Ford - Example

Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$



B-F equation gives:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node achieving **minimum is next-hop** in shortest path, used in forwarding table.



Distance Vector Algorithm

- ◆ $D_x(y)$ = estimate of least cost from x to y
 - x maintains distance vector $\mathbf{D}_x = [D_x(y): y \in N]$
- ◆ Node x :
 - knows cost to each neighbor v : $c(x,v)$
 - maintains its neighbors' distance vectors. For each neighbor v , x maintains:
 $\mathbf{D}_v = [D_v(y): y \in N]$



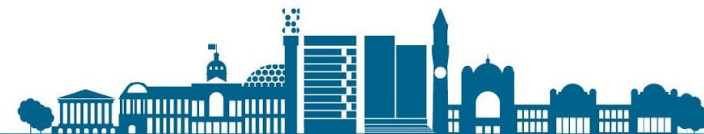
Distance Vector Algorithm

Key Idea:

- ◆ From time-to-time, each node sends its own distance vector estimate to neighbors
- ◆ When x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \text{ for each node } y \in N$$

- ◆ Under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$



Distance Vector Algorithm

Iterative, Asynchronous:


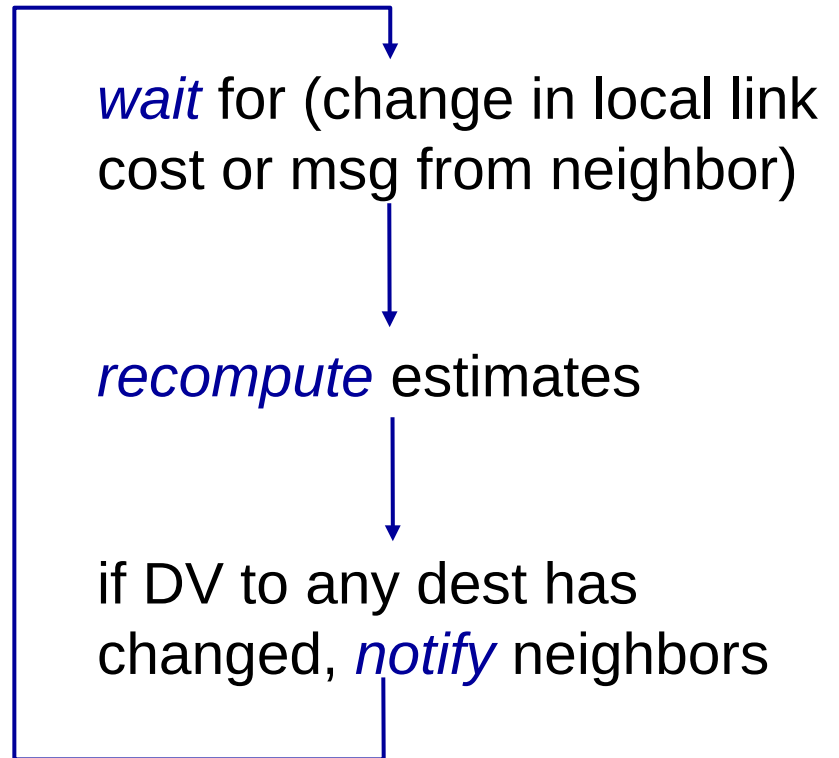
each local iteration
caused by:

- ◆ local link cost change
- ◆ DV update message from neighbor

Distributed:

- ◆ each node notifies neighbors *only* when its DV changes
- neighbors then notify their neighbors if necessary

Each Node:



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

**node x
table**

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

**node y
table**


		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

**node z
table**

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0



time



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x
table

from \ cost to			
	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

node y
table

from \ cost to			
	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

node z
table

from \ cost to			
	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

from \ cost to			
	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0



from \ cost to			
	x	y	z
x	0	2	7
y	2	0	1
z	7	1	0

from \ cost to			
	x	y	z
x	0	2	7
y	2	0	1
z	3	1	0

from \ cost to			
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0

from \ cost to			
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0


from \ cost to			
	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0



Distance Vector - Link Cost Changes

Link Cost Changes:

- ◆ Node detects local link cost change
- ◆ Updates routing info, recalculates distance vector
- ◆ If DV changes, notify neighbors




t_0 : y detects link-cost change, updates its DV, informs its neighbors.

t_1 : z receives update from y, updates its table, computes new least cost to x, sends its neighbors its DV.

t_2 : y receives z's update, updates its distance table. y's least costs do *not* change, so y does *not* send a message to z.

“Good news travels fast!”

What would happen if the link cost increases?



Comparison of LS and DV Algorithms

Message Complexity

- ◆ **LS**: with N nodes, E links, $O(NE)$ messages sent
- ◆ **DV**: exchange between neighbors only
 - convergence time varies

Speed of Convergence

- ◆ **LS**: $O(N^2)$ algorithm requires $O(NE)$ messages
 - may have oscillations
- ◆ **DV**: convergence time varies
 - may be routing loops
 - count-to-infinity problem


Robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect **link** cost
- each node computes only its own table

DV:

- DV node can advertise incorrect **path** cost
 - ▶ each node's table used by others
 - ▶ error propagate through the network



Hierarchical Routing

Our routing study thus far - idealization

- ◆ All routers identical
- ◆ Network is “flat” ... *not* true in practice

Scale: with 600 million destinations:

- ◆ Can't store all dest's in routing tables!
- ◆ Routing table exchange would swamp links!

Administrative Autonomy

- ◆ Internet = network of networks
- ◆ Each network admin may want to control routing in its own network



Hierarchical Routing

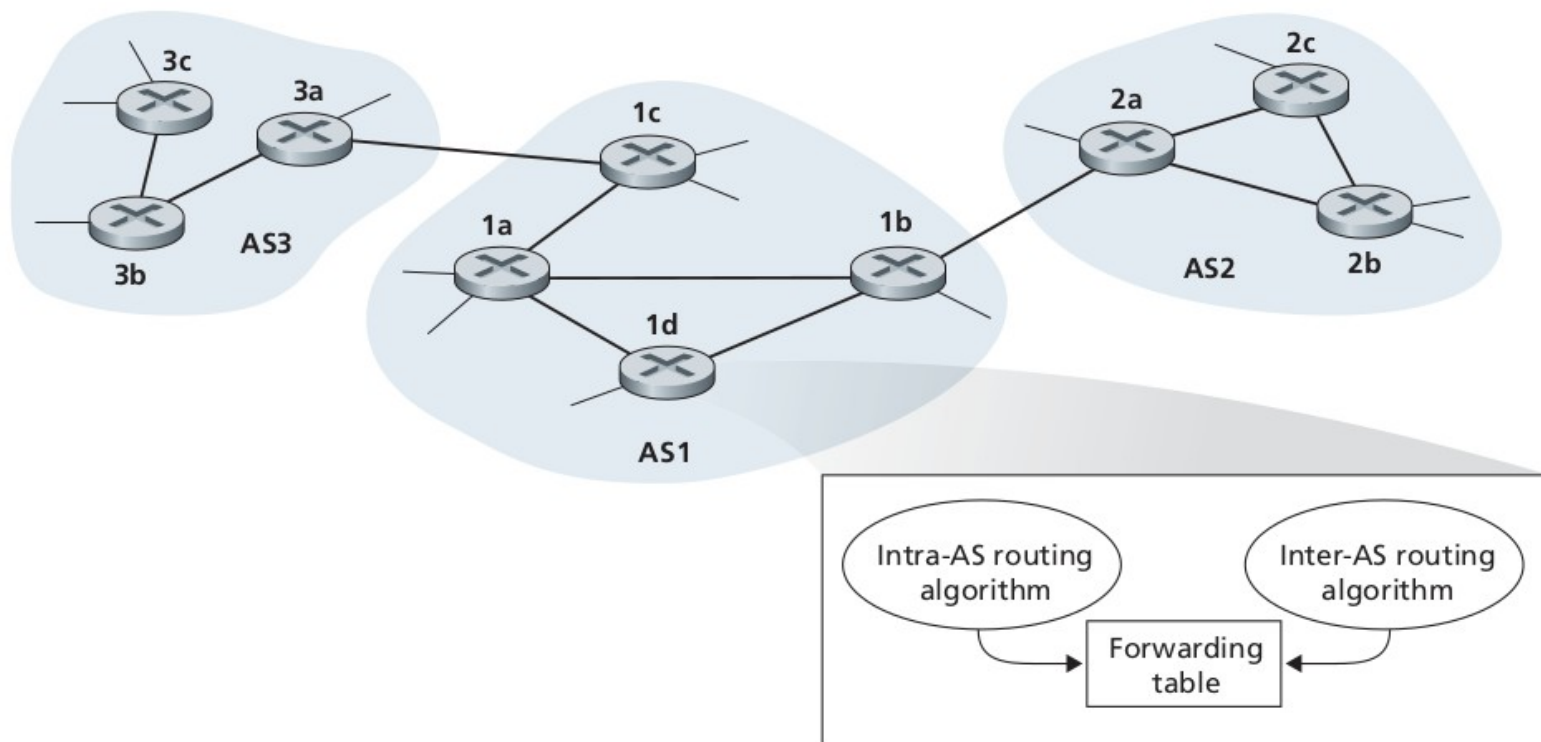
- ◆ Aggregate routers into regions, “Autonomous Systems” (AS)
- ◆ Routers in same AS run same routing protocol
 - “intra-AS” routing protocol
 - routers in different AS can run different intra-AS routing protocol

Gateway Router:


- ◆ At “edge” of its own AS
- ◆ Has link to router in another AS



Interconnected ASes



- ◆ Forwarding table configured by both intra- and inter-AS routing algorithm
 - **intra-AS** sets entries for internal destinations
 - **inter-AS & intra-AS** sets entries for external destinations




Inter-AS Tasks

- ◆ Suppose a router (1d) in AS1 receives datagram destined outside of AS1:
 - router should forward packet to gateway router, but which one?

AS1 must:


- 1) Learn which destds are reachable through AS2, which through AS3
- 2) Propagate this reachability info to all routers in AS1

Job of inter-AS routing!





Example: Setting Forwarding Table in Router 1d

- ◆ Suppose AS1 learns (via inter-AS protocol) that the subnet **x** is reachable via AS3 (gateway 1c), but not via AS2
 - inter-AS protocol propagates reachability info to all internal routers
- ◆ Router 1d determines from intra-AS routing info that its interface **l** is on the least cost path to 1c
 - installs forwarding table entry **(x,l)**



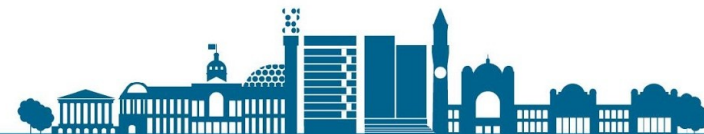
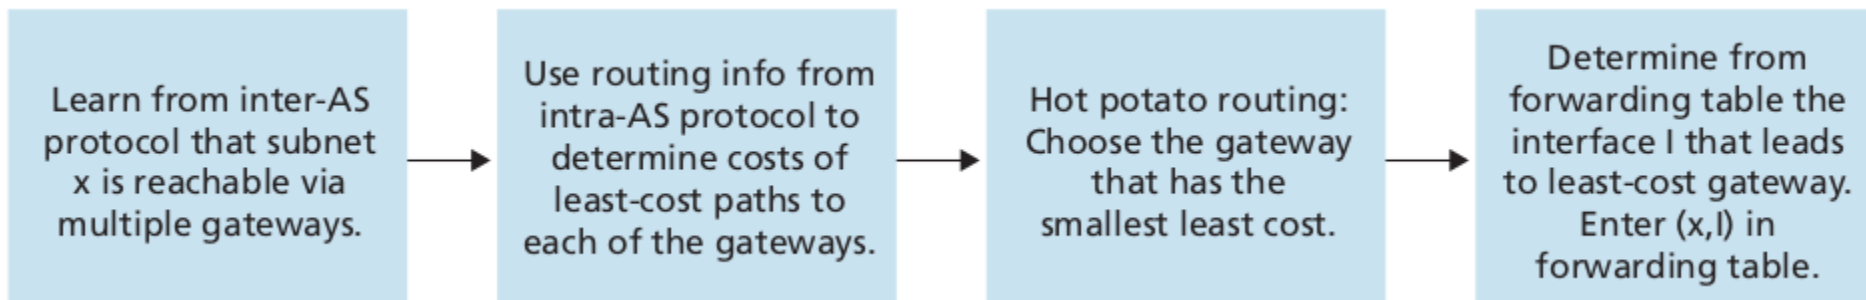
Example: Choosing Among Multiple ASes

- ◆ Now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ◆ To configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest **x**
 - This is also job of inter-AS routing protocol!



Example: Choosing Among Multiple ASes


- ◆ Now suppose AS1 learns from inter-AS protocol that subnet **x** is reachable from AS3 *and* from AS2.
- ◆ To configure forwarding table, router 1d must determine which gateway it should forward packets towards for dest **x**
 - This is also job of inter-AS routing protocol!
- ◆ *Hot Potato Routing*: send packet towards closest of two routers.



Summary

In this lecture, we have seen:

- ◆ The two main classes of routing algorithms i.e. Link State and Distance Vector routing.
- ◆ Dijkstra's Algorithm is a Link State algorithm.
- ◆ Bellman-Ford as a Distance Vector algorithm.
- ◆ The importance of Hierarchical Routing.



References / Links

- ◆ Chapter #5: **The Network Layer: Control Plane**,
Computer Networking: A Top-Down Approach (7th edition)
by Kurose & Ross

